

# **CMSC 474, Game Theory**

## **4b. Game-Tree Search**

Dana Nau

University of Maryland

# Finite perfect-information zero-sum games

- **Finite:**

- finitely many agents, actions, states, histories

- **Perfect information:**

- Every agent knows
  - all of the players' utility functions
  - all of the players' actions and what they do
  - the history and current state
- No simultaneous actions – agents move one-at-a-time

- **Constant sum (or zero-sum):**

- Constant  $k$  such that regardless of how the game ends,
  - $\sum_{i=1,\dots,n} u_i = k$
- For every such game, there's an equivalent game in which  $k = 0$

# Examples

- **Deterministic:**

- outcomes depend only on the players' moves
- tic-tac-toe, qubic, connect-four, mancala, reversi (othello), chess, checkers, go

- **Stochastic:**

- outcomes depend partly on chance (e.g., dice rolls)
- backgammon, monopoly, yahtzee, parcheesi, roulette, craps

- For now, we'll consider just the deterministic games


# Outline

- Non-credible threats
- Restatement of the Minimax Theorem
- Game trees
- The minimax algorithm
- $\alpha$ - $\beta$  pruning
- Resource limits, approximate evaluation
  
- Most of this isn't in the game-theory book
  - I'll post some material on Piazza
- For further information, look at one of the following
  - Russell & Norvig's *Artificial Intelligence: A Modern Approach*
    - Chapter 6 of the 2<sup>nd</sup> edition
    - Chapter 5 of the 3<sup>rd</sup> edition

# No Non-Credible Threats



- Recall this example from Chapter 4:
- If the agents can announce their strategies beforehand, agent 1 might want to announce (B,H)
- Non-credible threat:
  - Agent 1 is trying to keep 2 from choosing F, by threatening to retaliate with H
  - But if 2 chooses F anyway, then it wouldn't be rational for 1 to do H, because H also hurts 1
- In zero-sum games, non-credible threats can't occur
  - In a zero-sum game, if H hurts 2 then it helps 1



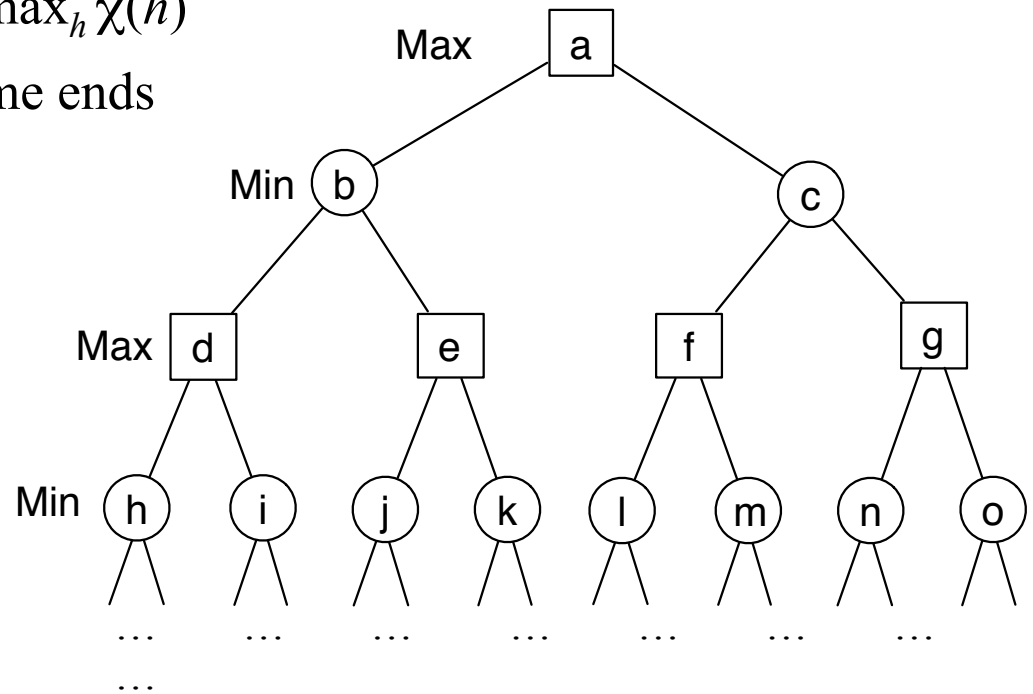
# Minimax Theorem (Restated)

- Instead of  $u_1$  and  $u_2$ , write  $u$  and  $-u$ 
  - Call player 1 Max (wants to maximize  $u$ ), call Max's strategy  $s$
  - Call player 2 Min (wants to minimize  $u$ ), call Min's strategy  $t$
- **Theorem.** Let  $G$  be any finite two-player zero-sum game. Then
  - Max's expected utility in any Nash equilibrium  
= Max's maxmin value = Max's minmax value
  - In other words, for every Nash equilibrium  $(s^*, t^*)$ ,
    - $u(s^*, t^*) = \min_s \max_t u(s, t) = \max_s \min_t u(s, t)$
- **Corollary.** There are pure strategies  $s^*$  and  $t^*$  that satisfy the theorem.
- Terminology
  - $u(s^*, t^*)$  is called the *minimax value* of  $G$
  - $s^*$  and  $t^*$  are sometimes called *optimal strategies* or *perfect play*

# Terminology

- **Root node**: where the game starts
- **Max (or Min) node**:
  - a node where it's Max's (or Min's) move
  - usually draw Max nodes as squares, Min nodes as circles
- A node's **children**:  $\{\sigma(h, a) \mid a \in \chi(h)\}$ 
  - **Branching factor**  $b = \max_h \chi(h)$
- **Terminal nodes**: where game ends

$H = \{\text{nonterminal nodes}\}$   
 $Z = \{\text{terminal nodes}\}$   
 $\rho(h) = \text{the player to move at } h$   
 $\chi(h) = \{\text{available actions at } h\}$   
 $\sigma(h, a) = \text{child of } h \text{ produced by } a$   
 $\mathbf{u}(h) = \text{utility profile for } h$   
 $\mathbf{v}[i] = i\text{'th element of } \mathbf{v}$



# Minimax Algorithm

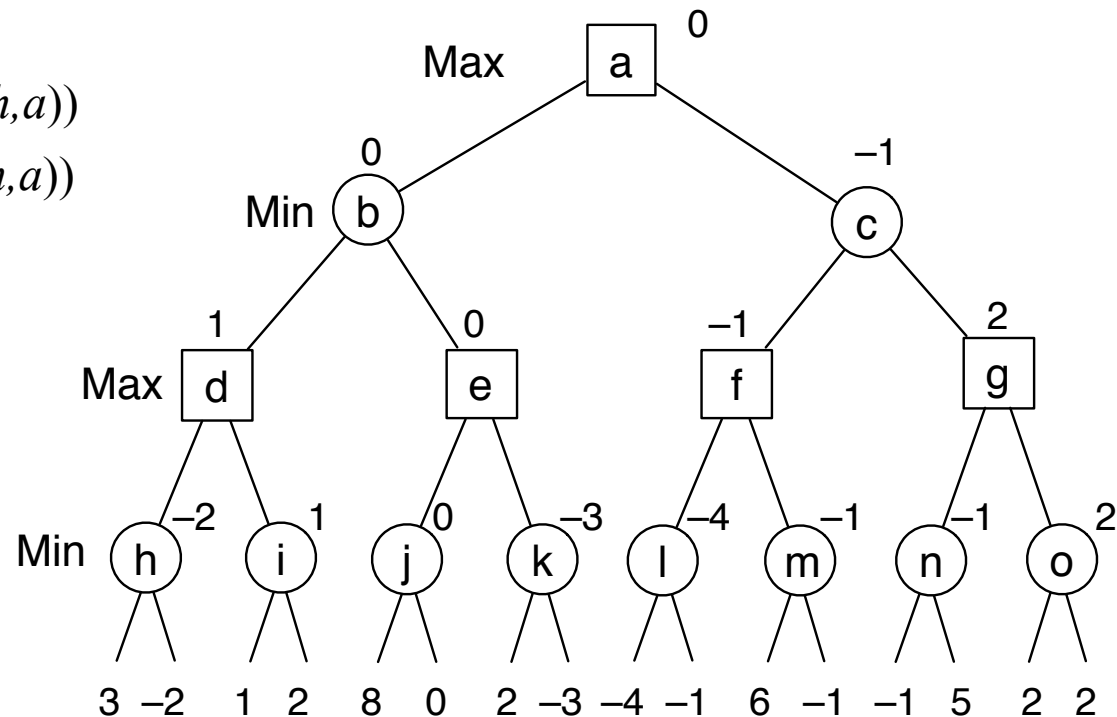
- Backward induction (Chapter 4)
  - Simplified for 2-player zero-sum games

function Minimax( $h$ )  
 if  $h \in Z$  then return  $u(h)$   
 else  $\rho(h) = \text{Max}$  then  
     return  $\max_{a \in \chi(h)} \text{Minimax}(\sigma(h, a))$   
 else return  $\min_{a \in \chi(h)} \text{Minimax}(\sigma(h, a))$

$Z = \{\text{terminal nodes}\}$   
 $\rho(h)$  = the player to move at  $h$   
 $\chi(h) = \{\text{available actions at } h\}$   
 $\sigma(h, a) = \text{child of } h \text{ produced by } a$

function Minimax-choice( $h$ )  
 if  $h \in Z$  then return *error*  
 else if  $\rho(h) = \text{Max}$  then  
     return  $\arg \max_{a \in \chi(h)} \text{Minimax}(\sigma(h, a))$   
 else return  $\arg \min_{a \in \chi(h)} \text{Minimax}(\sigma(h, a))$

**Poll 4.2:** what would Minimax-choice return at nodes b and c?





# Complexity Analysis

- Let  $b$  = branching factor,  $D$  = height of tree
- Space complexity  
= (max path length)  $\times$  (space needed to store the path)  
=  $O(bD)$
- Time complexity = size of the game tree =  $O(b^D)$
- Example: chess
  - for “reasonable” chess games,  $b \approx 35$ ,  $D \approx 100$
  - $b^h \approx 35^{100} \approx 10^{135}$  nodes
- $\approx 10^{87}$  particles in the universe
  - $10^{135}$  nodes is  $\approx 10^{48}$  times the number of particles in the universe

# Limited-Depth Minimax

- Upper bound  $d$  on search depth
  - Running time  $O(b^d)$
  - Space  $O(bd)$
- If  $d \geq \text{height}(h)$ , get  $u(h)$ 
  - Otherwise, approximation
- **static evaluation function**  $e(h)$ 
  - returns estimate of  $u(h)$

$H = \{\text{nonterminal nodes}\}$   
 $Z = \{\text{terminal nodes}\}$   
 $\rho(h) = \text{the player to move at } h$   
 $\chi(h) = \{\text{available actions at } h\}$   
 $\sigma(h, a) = \text{child of } h \text{ produced by } a$   
 $\mathbf{u}(h) = \text{utility profile for } h$   
 $\mathbf{v}[i] = i\text{'th element of } \mathbf{v}$

```

function LD-minimax( $h, d$ )
    if  $h \in Z$  then return  $u(h)$ 
    else if  $d = 0$  then return  $e(h)$ 
    else if  $\rho(h) = \text{Max}$  then
        return  $\max_{a \in \chi(h)} \text{LD-minimax}(\sigma(h, a), d-1)$ 
    else return  $\min_{a \in \chi(h)} \text{LD-minimax}(\sigma(h, a), d-1)$ 
    
```


```

function LD-minimax-choice( $h, d$ )
    if  $h \in Z$  or  $d = 0$  then return error
    else if  $\rho(h) = \text{Max}$  then
        return  $\arg \max_{a \in \chi(h)} \text{LD-minimax}(\sigma(h, a), d-1)$ 
    else return  $\arg \min_{a \in \chi(h)} \text{LD-minimax}(\sigma(h, a), d-1)$ 
    
```


- LD-minimax-choice returns an action
- Call it again at every move
  - Why?

# Evaluation Functions

- $e(h)$  is often a weighted sum of features
  - $e(h) = c_1 f_1(h) + c_2 f_2(h) + \dots + c_n f_n(h)$
- E.g., chess
  - $c_1 \times \text{material} + c_2 \times \text{mobility} + c_3 \times \text{king safety} + c_4 \times \text{center control} + \dots$ 
    - $\text{material} = 1 \times (\text{your pawns} - \text{opponent's pawns}) + 3 \times (\text{your knights} - \text{opponent's knights}) + \dots$



Black to move  
White slightly better





White to move  
Black winning

# Exact Values for $e$ Don't Matter

MAX


MIN



- Behavior is preserved under any **monotonic** transformation of  $e$ 
  - Only the order matters

# Pruning

- Minimax and LD-minimax examine nodes that don't need to be examined




```

function LD-minimax( $h, d$ )
  if  $h \in Z$  then return  $u(h)$ 
  else if  $d = 0$  then return  $e(h)$ 
  else if  $\rho(h) = \text{Max}$  then
     $v^* \leftarrow -\infty$ 
    for every  $a \in \chi(h)$  do
       $v \leftarrow \text{LD-minimax}(\sigma(h, a), d-1)$ 
      if  $v^* < v$  then  $v^* \leftarrow v$ 
  else
     $v^* \leftarrow \infty$ 
    for every  $a \in \chi(h)$  do
       $v \leftarrow \text{LD-minimax}(\sigma(h, a), d-1)$ 
      if  $v^* > v$  then  $v^* \leftarrow v$ 
  return  $v$ 
  
```

# Pruning

- For Max,  $b$  is better than  $f$ 
  - If Max is rational then Max will never choose  $f$
  - Don't need to examine any more nodes below  $f$




```

function LD-minimax( $h, d$ )
    if  $h \in Z$  then return  $u(h)$ 
    else if  $d = 0$  then return  $e(h)$ 
    else if  $\rho(h) = \text{Max}$  then
         $v^* \leftarrow -\infty$ 
        for every  $a \in \chi(h)$  do
             $v \leftarrow \text{LD-minimax}(\sigma(h, a), d-1)$ 
            if  $v^* < v$  then  $v^* \leftarrow v$ 
    else
         $v^* \leftarrow \infty$ 
        for every  $a \in \chi(h)$  do
             $v \leftarrow \text{LD-minimax}(\sigma(h, a), d-1)$ 
            if  $v^* > v$  then  $v^* \leftarrow v$ 
    return  $v$ 
    
```


# Pruning

- Don't know whether  $h$  is better or worse than  $b$



# Pruning


- Still don't know whether  $h$  is better or worse than  $b$






# Pruning

- $h$  is worse than  $b$ 
  - Don't need to examine any more nodes below  $h$
- $v(a) = 3$




# Alpha Cutoff

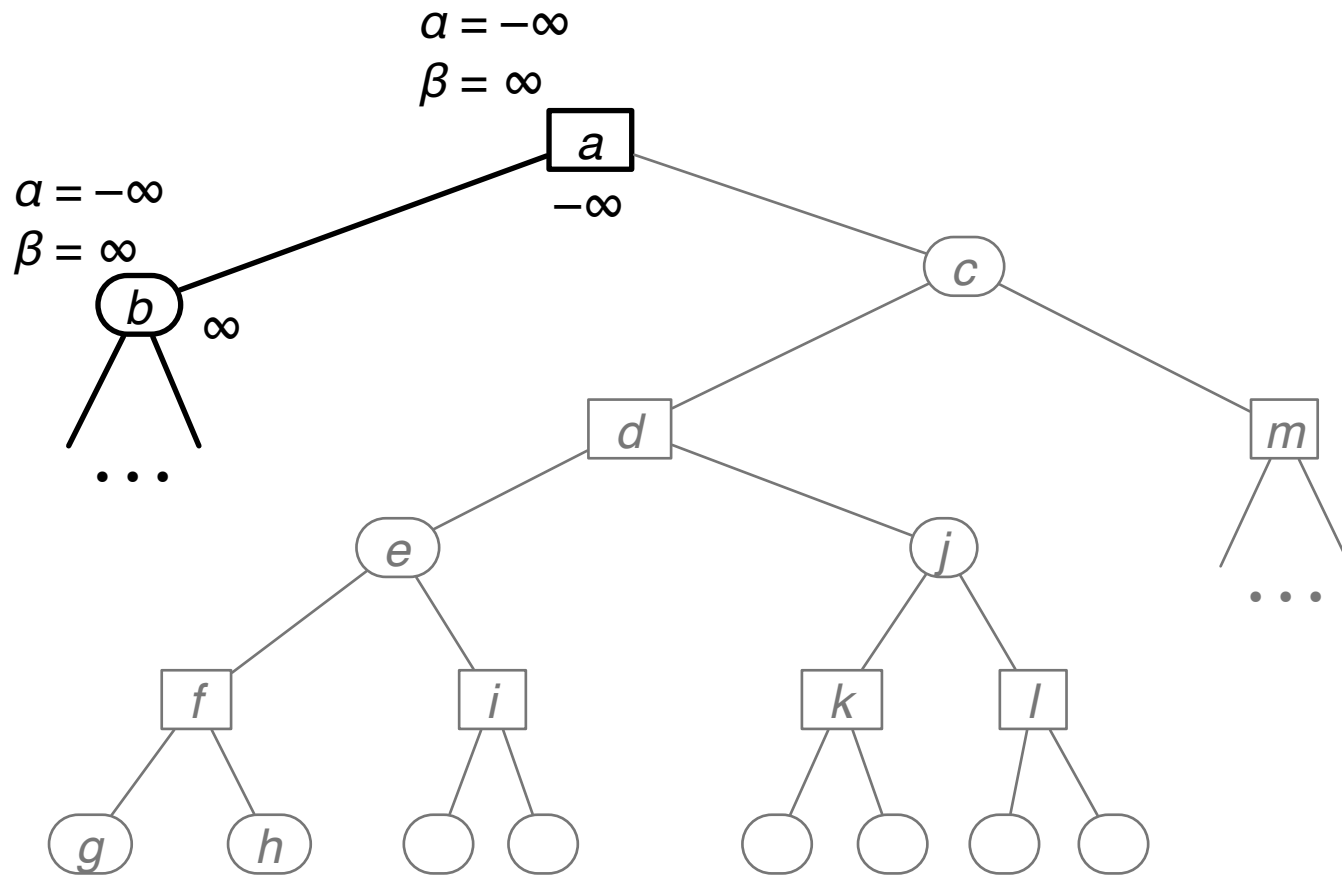
- At Min nodes, let  $\alpha$  = largest lower bound on ancestors
  - At  $s$ ,  $\alpha = \max(a, b, c)$
- Suppose  $d < \alpha$ 
  - If the game ever reaches  $s$ , Max's payoff will be  $< \alpha$
- To get payoff  $\geq \alpha$ ,
  - Max will move elsewhere at  $p$ ,  $q$ , or  $r$
  - The game won't ever reach  $s$
- What if  $d = \alpha$ ?



# Beta Cutoff

- At Max nodes, let  $\beta$  = smallest upper bound on ancestors
  - At  $s$ ,  $\beta = \min(a, b, c)$
- Suppose  $d > \beta$ 
  - If the game ever reaches  $s$ , Max's payoff will be  $> \beta$
- To make Max's payoff  $\leq \beta$ 
  - Min will move elsewhere at  $p$ ,  $q$ , or  $r$
  - The game won't ever reach  $s$
- What if  $d = \beta$ ?






```

function Alpha-Beta( $h, d, \alpha, \beta$ )
  if  $h \in Z$  then return  $u(h)$ 
  else if  $d = 0$  then return  $e(h)$ 
  else if  $\rho(h) = \text{Max}$  then
     $v \leftarrow -\infty$ 
    for every  $a \in \chi(h)$  do
       $v \leftarrow \max(v, \text{Alpha-Beta}(\sigma(h, a), d-1, \alpha, \beta))$ 
      if  $v \geq \beta$  then return  $v$  //  $\beta$  cutoff
      else  $\alpha \leftarrow \max(\alpha, v)$ 
    return  $v$ 
  else
     $v \leftarrow \infty$ 
    for every  $a \in \chi(h)$  do
       $v \leftarrow \min(v, \text{Alpha-Beta}(\sigma(h, a), d-1, \alpha, \beta))$ 
      if  $v \leq \alpha$  then return  $v$  //  $\alpha$  cutoff
      else  $\beta \leftarrow \min(\beta, v)$ 
    return  $v$ 

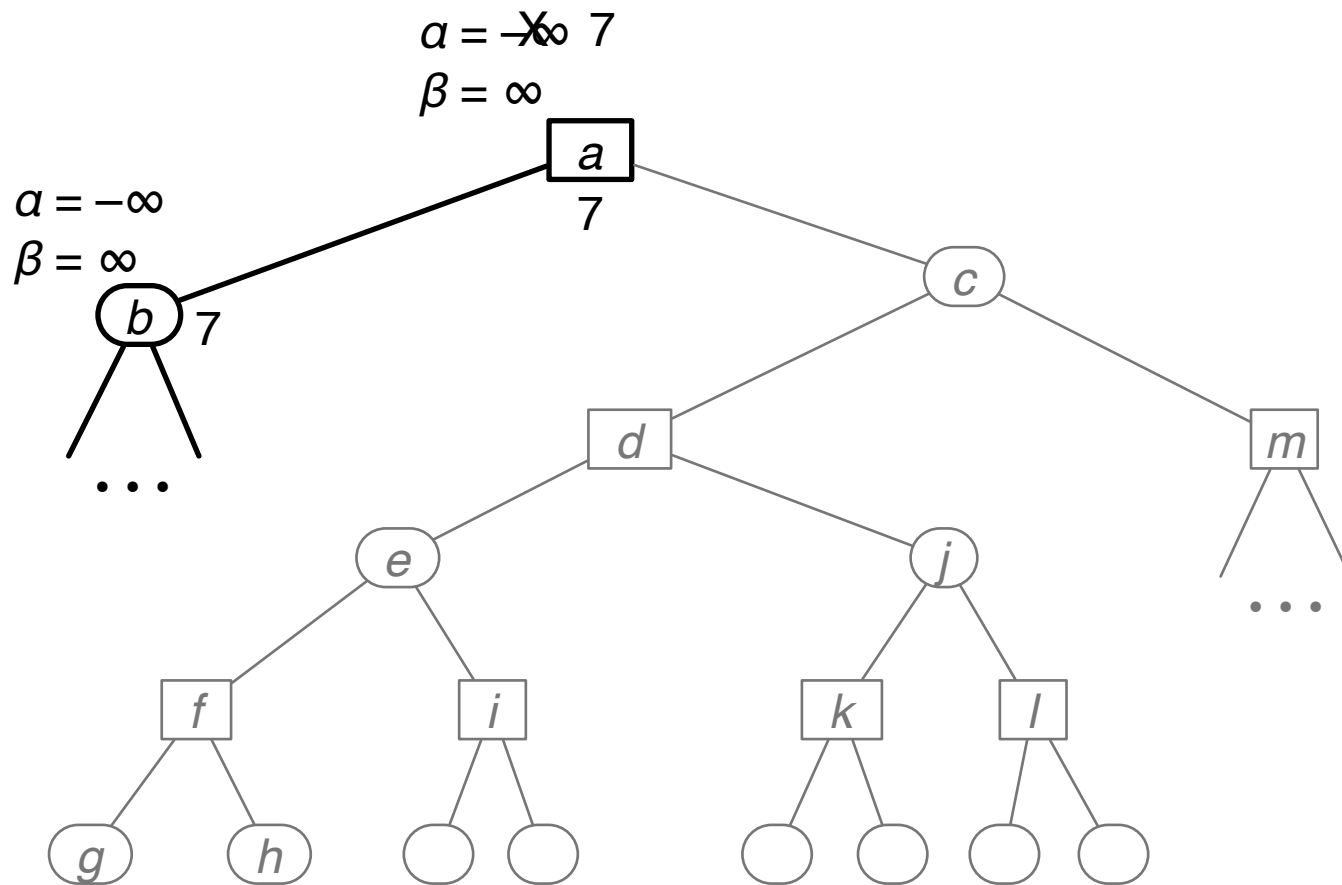
```



```

function Alpha-Beta( $h, d, \alpha, \beta$ )
  if  $h \in Z$  then return  $u(h)$ 
  else if  $d = 0$  then return  $e(h)$ 
  else if  $\rho(h) = \text{Max}$  then
     $v \leftarrow -\infty$ 
    for every  $a \in \chi(h)$  do
       $v \leftarrow \max(v, \text{Alpha-Beta}(\sigma(h, a), d-1, \alpha, \beta))$ 
      if  $v \geq \beta$  then return  $v$  //  $\beta$  cutoff
      else  $\alpha \leftarrow \max(\alpha, v)$ 
    return  $v$ 
  else
     $v \leftarrow \infty$ 
    for every  $a \in \chi(h)$  do
       $v \leftarrow \min(v, \text{Alpha-Beta}(\sigma(h, a), d-1, \alpha, \beta))$ 
      if  $v \leq \alpha$  then return  $v$  //  $\alpha$  cutoff
      else  $\beta \leftarrow \min(\beta, v)$ 
    return  $v$ 


```



```

function Alpha-Beta( $h, d, \alpha, \beta$ )
  if  $h \in Z$  then return  $u(h)$ 
  else if  $d = 0$  then return  $e(h)$ 
  else if  $\rho(h) = \text{Max}$  then
     $v \leftarrow -\infty$ 
    for every  $a \in \chi(h)$  do
       $v \leftarrow \max(v, \text{Alpha-Beta}(\sigma(h, a), d-1, \alpha, \beta))$ 
      if  $v \geq \beta$  then return  $v$  //  $\beta$  cutoff
      else  $\alpha \leftarrow \max(\alpha, v)$ 
    return  $v$ 
  else
     $v \leftarrow \infty$ 
    for every  $a \in \chi(h)$  do
       $v \leftarrow \min(v, \text{Alpha-Beta}(\sigma(h, a), d-1, \alpha, \beta))$ 
      if  $v \leq \alpha$  then return  $v$  //  $\alpha$  cutoff
      else  $\beta \leftarrow \min(\beta, v)$ 
    return  $v$ 

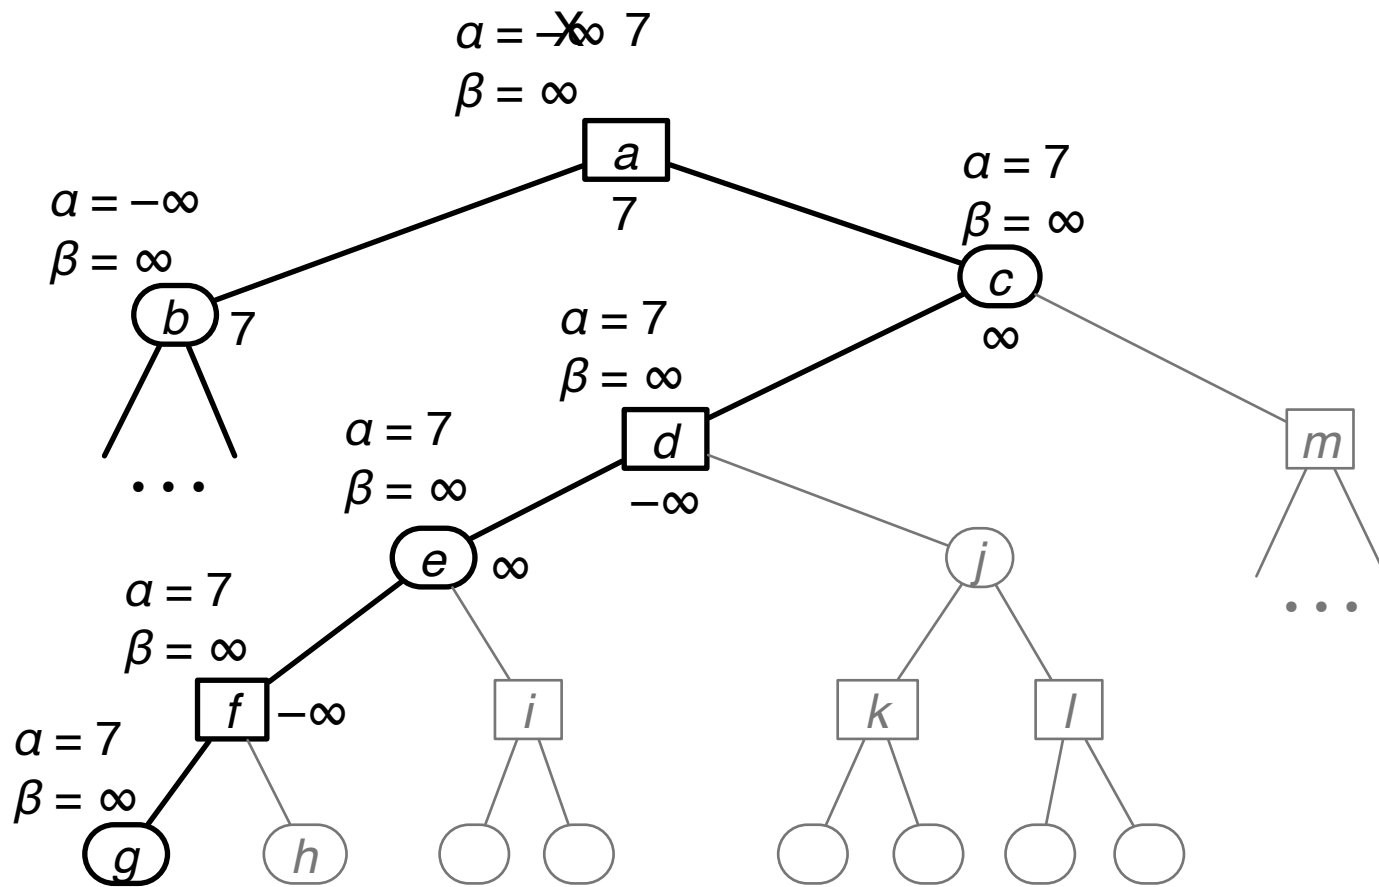
```



```

function Alpha-Beta( $h, d, \alpha, \beta$ )
  if  $h \in Z$  then return  $u(h)$ 
  else if  $d = 0$  then return  $e(h)$ 
  else if  $\rho(h) = \text{Max}$  then
     $v \leftarrow -\infty$ 
    for every  $a \in \chi(h)$  do
       $v \leftarrow \max(v, \text{Alpha-Beta}(\sigma(h, a), d-1, \alpha, \beta))$ 
      if  $v \geq \beta$  then return  $v$  //  $\beta$  cutoff
      else  $\alpha \leftarrow \max(\alpha, v)$ 
    return  $v$ 
  else
     $v \leftarrow \infty$ 
    for every  $a \in \chi(h)$  do
       $v \leftarrow \min(v, \text{Alpha-Beta}(\sigma(h, a), d-1, \alpha, \beta))$ 
      if  $v \leq \alpha$  then return  $v$  //  $\alpha$  cutoff
      else  $\beta \leftarrow \min(\beta, v)$ 
    return  $v$ 

```



function Alpha-Beta( $h, d, \alpha, \beta$ )

if  $h \in Z$  then return  $u(h)$

else if  $d = 0$  then return  $e(h)$

else if  $\rho(h) = \text{Max}$  then

$v \leftarrow -\infty$

for every  $a \in \chi(h)$  do

$v \leftarrow \max(v, \text{Alpha-Beta}(\sigma(h, a), d-1, \alpha, \beta))$

if  $v \geq \beta$  then return  $v$  //  $\beta$  cutoff

else  $\alpha \leftarrow \max(\alpha, v)$

return  $v$

else

$v \leftarrow \infty$

for every  $a \in \chi(h)$  do

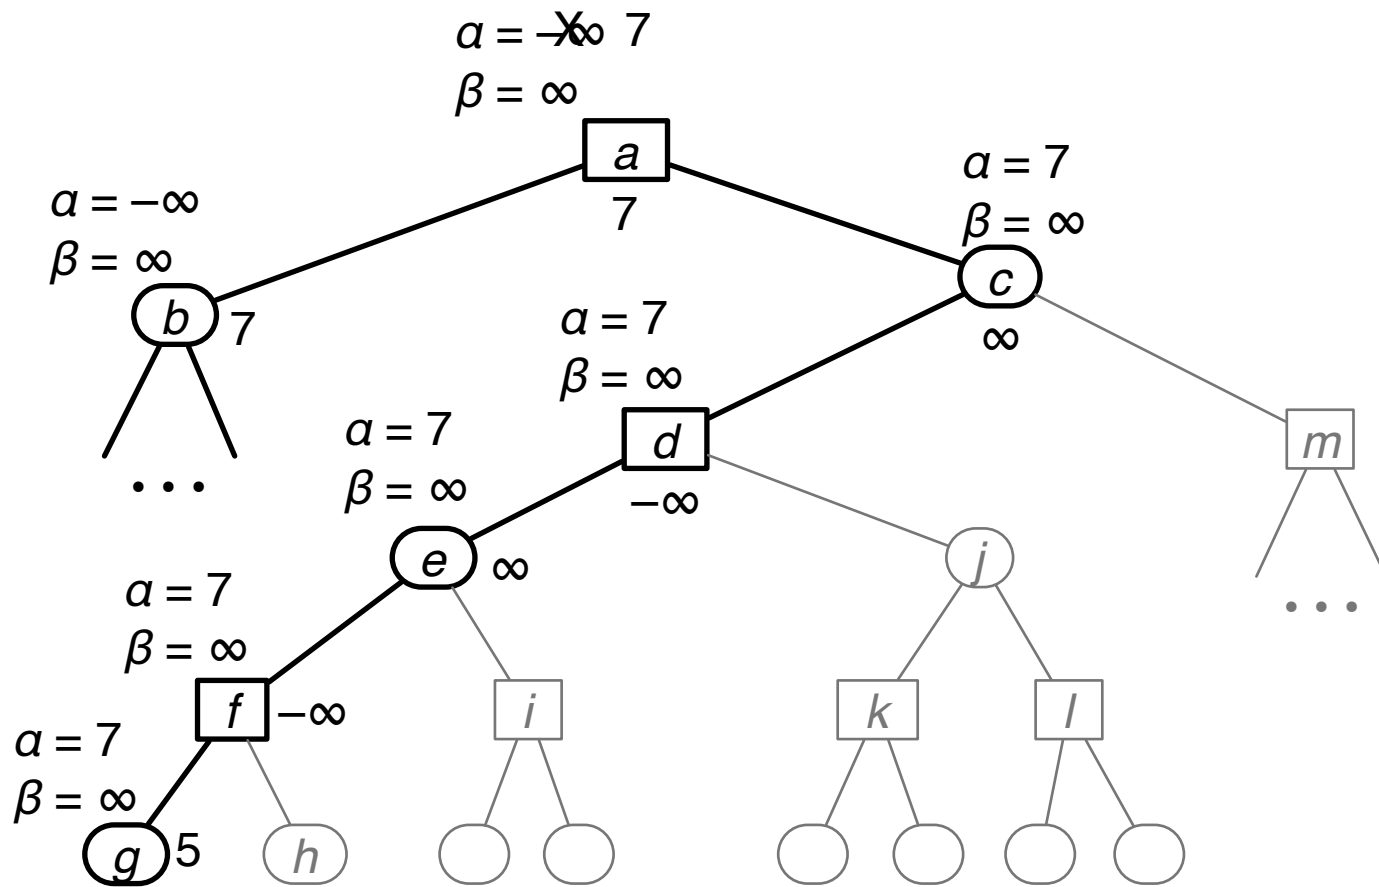
$v \leftarrow \min(v, \text{Alpha-Beta}(\sigma(h, a), d-1, \alpha, \beta))$

if  $v \leq \alpha$  then return  $v$  //  $\alpha$  cutoff

else  $\beta \leftarrow \min(\beta, v)$

return  $v$





function Alpha-Beta( $h, d, \alpha, \beta$ )

if  $h \in Z$  then return  $u(h)$

else if  $d = 0$  then return  $e(h)$

else if  $\rho(h) = \text{Max}$  then

$v \leftarrow -\infty$

for every  $a \in \chi(h)$  do

$v \leftarrow \max(v, \text{Alpha-Beta}(\sigma(h, a), d-1, \alpha, \beta))$

if  $v \geq \beta$  then return  $v$  //  $\beta$  cutoff

else  $\alpha \leftarrow \max(\alpha, v)$

return  $v$

else

$v \leftarrow \infty$

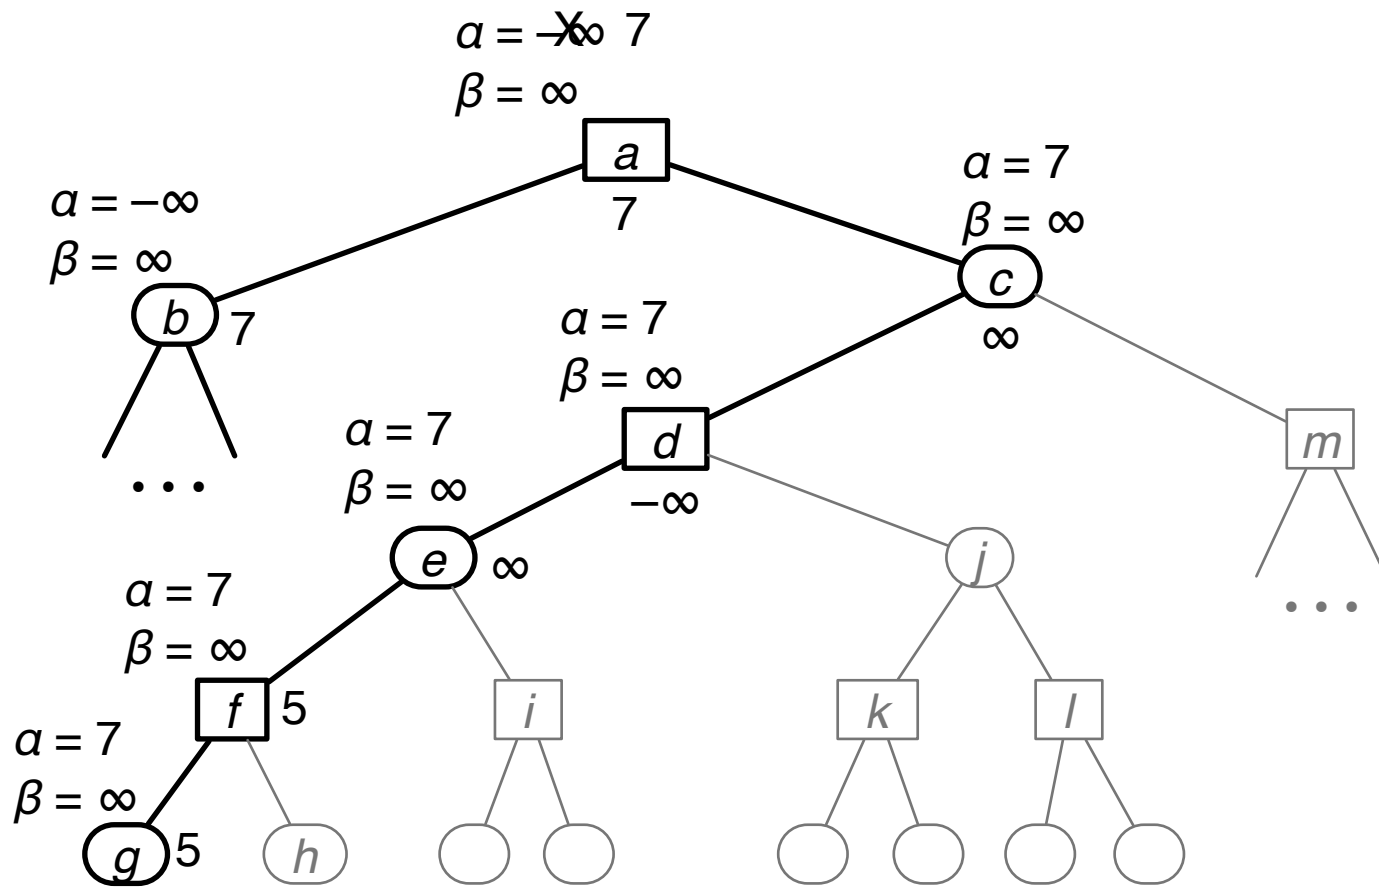
for every  $a \in \chi(h)$  do

$v \leftarrow \min(v, \text{Alpha-Beta}(\sigma(h, a), d-1, \alpha, \beta))$

if  $v \leq \alpha$  then return  $v$  //  $\alpha$  cutoff

else  $\beta \leftarrow \min(\beta, v)$

return  $v$



function Alpha-Beta( $h, d, \alpha, \beta$ )

if  $h \in Z$  then return  $u(h)$

else if  $d = 0$  then return  $e(h)$

else if  $\rho(h) = \text{Max}$  then

$v \leftarrow -\infty$

for every  $a \in \chi(h)$  do

$v \leftarrow \max(v, \text{Alpha-Beta}(\sigma(h, a), d-1, \alpha, \beta))$

if  $v \geq \beta$  then return  $v$  //  $\beta$  cutoff

else  $\alpha \leftarrow \max(\alpha, v)$

return  $v$

else

$v \leftarrow \infty$

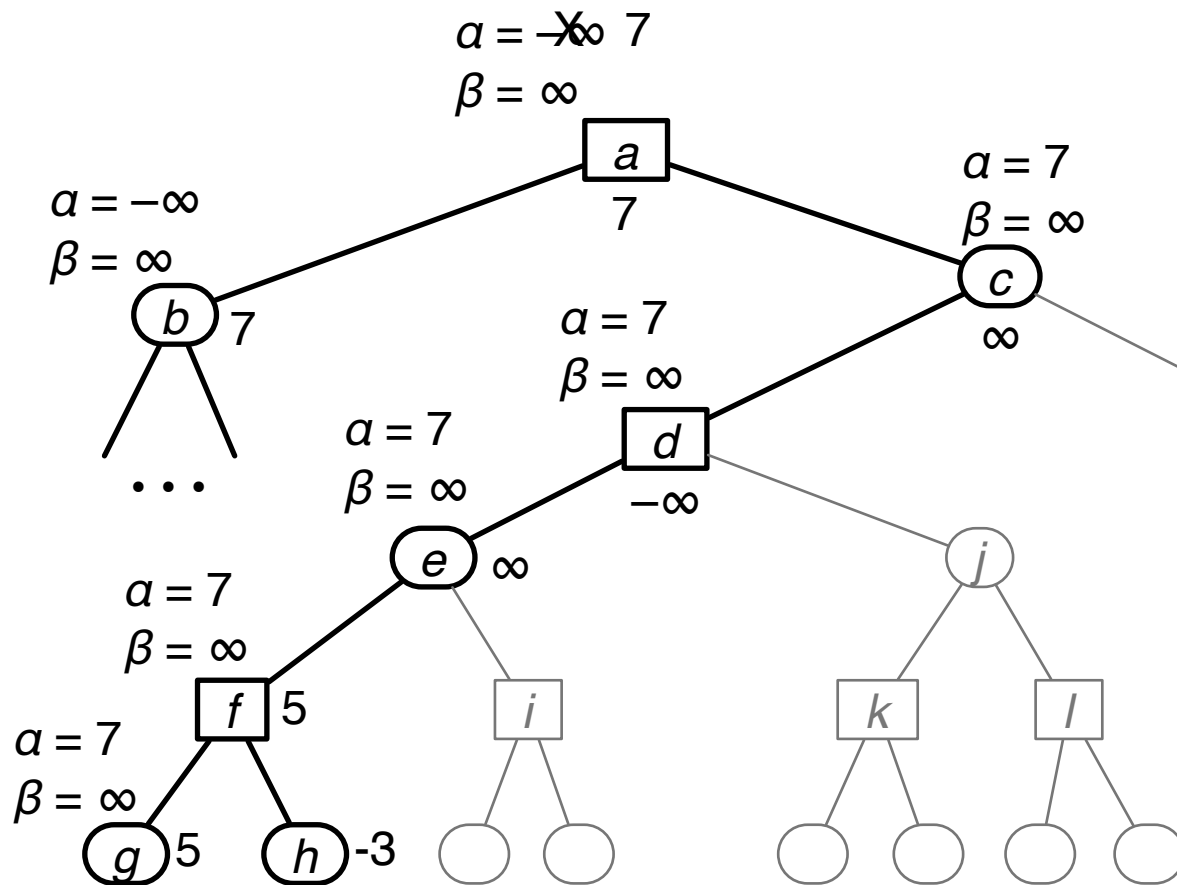
for every  $a \in \chi(h)$  do

$v \leftarrow \min(v, \text{Alpha-Beta}(\sigma(h, a), d-1, \alpha, \beta))$

if  $v \leq \alpha$  then return  $v$  //  $\alpha$  cutoff

else  $\beta \leftarrow \min(\beta, v)$

return  $v$



function Alpha-Beta( $h, d, \alpha, \beta$ )

if  $h \in Z$  then return  $u(h)$

else if  $d = 0$  then return  $e(h)$

else if  $\rho(h) = \text{Max}$  then

$v \leftarrow -\infty$

for every  $a \in \chi(h)$  do

$v \leftarrow \max(v, \text{Alpha-Beta}(\sigma(h, a), d-1, \alpha, \beta))$

if  $v \geq \beta$  then return  $v$  //  $\beta$  cutoff

else  $\alpha \leftarrow \max(\alpha, v)$

return  $v$

else

$v \leftarrow \infty$

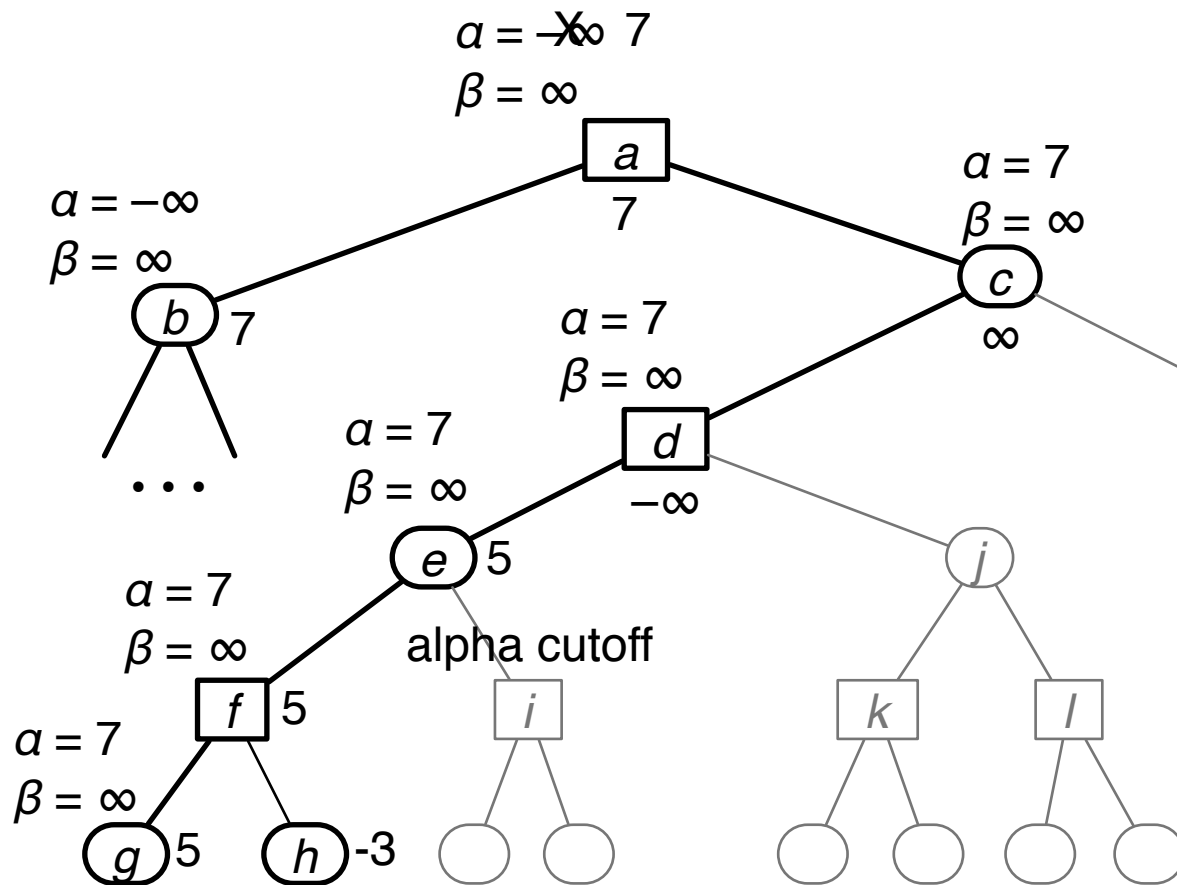
for every  $a \in \chi(h)$  do

$v \leftarrow \min(v, \text{Alpha-Beta}(\sigma(h, a), d-1, \alpha, \beta))$

if  $v \leq \alpha$  then return  $v$  //  $\alpha$  cutoff

else  $\beta \leftarrow \min(\beta, v)$

return  $v$



function Alpha-Beta( $h, d, \alpha, \beta$ )

if  $h \in Z$  then return  $u(h)$

else if  $d = 0$  then return  $e(h)$

else if  $\rho(h) = \text{Max}$  then

$v \leftarrow -\infty$

for every  $a \in \chi(h)$  do

$v \leftarrow \max(v, \text{Alpha-Beta}(\sigma(h, a), d-1, \alpha, \beta))$

if  $v \geq \beta$  then return  $v$  //  $\beta$  cutoff

else  $\alpha \leftarrow \max(\alpha, v)$

return  $v$

else

$v \leftarrow \infty$

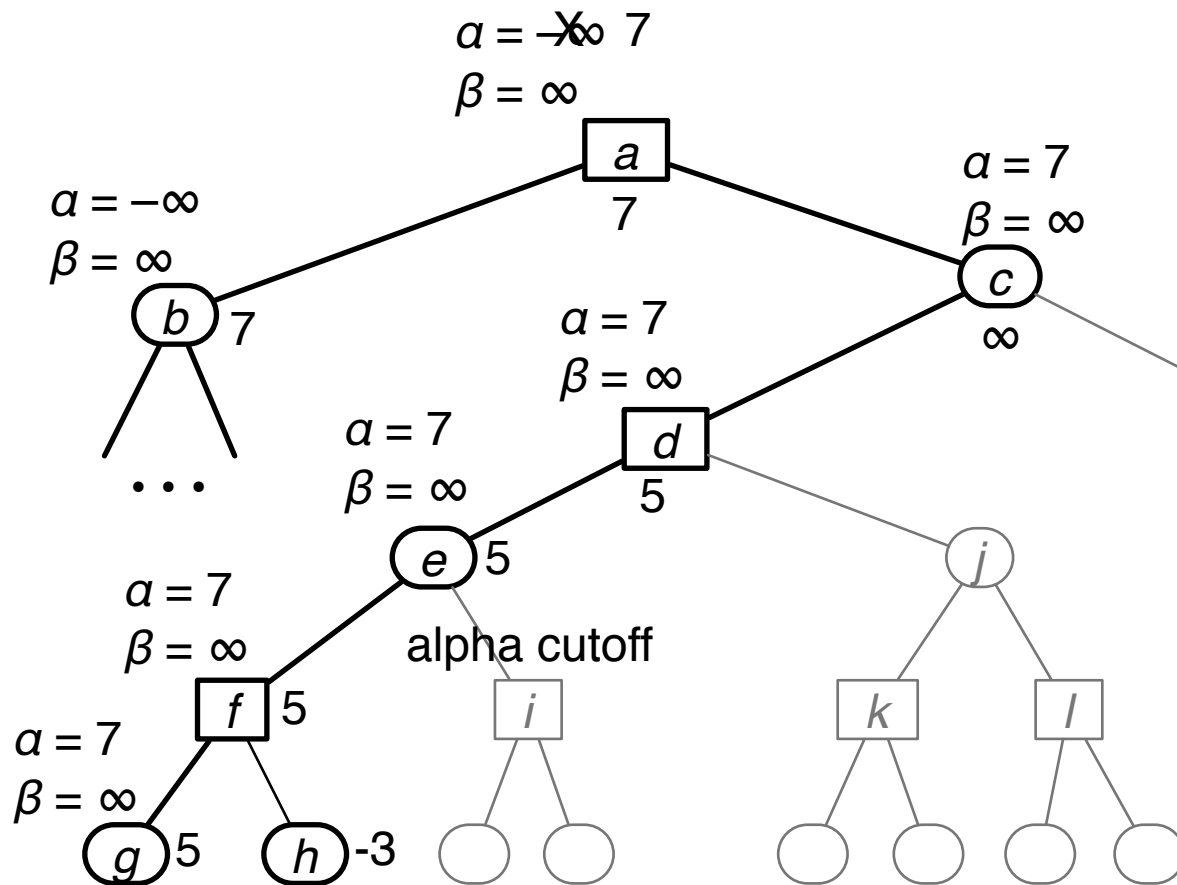
for every  $a \in \chi(h)$  do

$v \leftarrow \min(v, \text{Alpha-Beta}(\sigma(h, a), d-1, \alpha, \beta))$

if  $v \leq \alpha$  then return  $v$  //  $\alpha$  cutoff

else  $\beta \leftarrow \min(\beta, v)$

return  $v$



function Alpha-Beta( $h, d, \alpha, \beta$ )

if  $h \in Z$  then return  $u(h)$

else if  $d = 0$  then return  $e(h)$

else if  $\rho(h) = \text{Max}$  then

$v \leftarrow -\infty$

for every  $a \in \chi(h)$  do

$v \leftarrow \max(v, \text{Alpha-Beta}(\sigma(h, a), d-1, \alpha, \beta))$

if  $v \geq \beta$  then return  $v$  //  $\beta$  cutoff

else  $\alpha \leftarrow \max(\alpha, v)$

return  $v$

else

$v \leftarrow \infty$

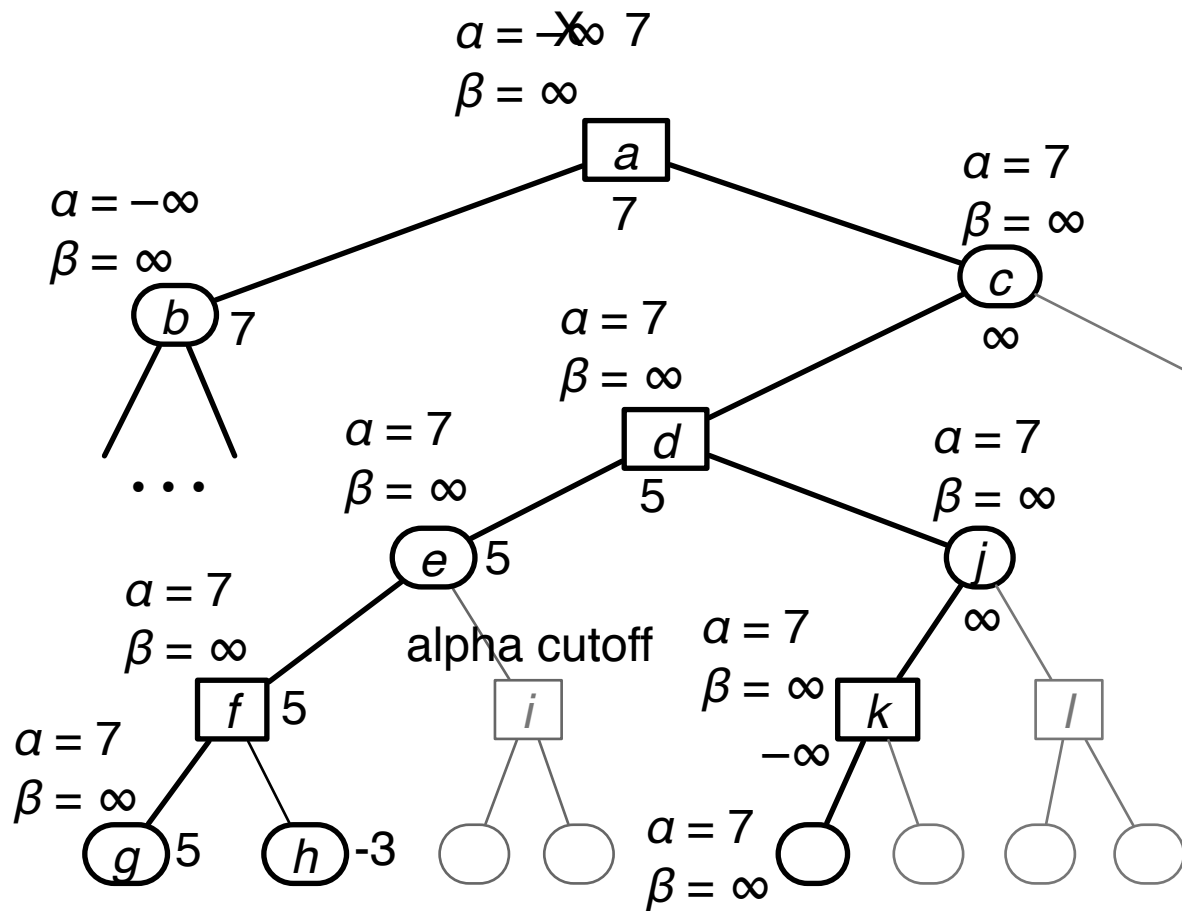
for every  $a \in \chi(h)$  do

$v \leftarrow \min(v, \text{Alpha-Beta}(\sigma(h, a), d-1, \alpha, \beta))$

if  $v \leq \alpha$  then return  $v$  //  $\alpha$  cutoff

else  $\beta \leftarrow \min(\beta, v)$

return  $v$



function Alpha-Beta( $h, d, \alpha, \beta$ )

if  $h \in Z$  then return  $u(h)$

else if  $d = 0$  then return  $e(h)$

else if  $\rho(h) = \text{Max}$  then

$v \leftarrow -\infty$

for every  $a \in \chi(h)$  do

$v \leftarrow \max(v, \text{Alpha-Beta}(\sigma(h, a), d-1, \alpha, \beta))$

if  $v \geq \beta$  then return  $v$  //  $\beta$  cutoff

else  $\alpha \leftarrow \max(\alpha, v)$

return  $v$

else

$v \leftarrow \infty$

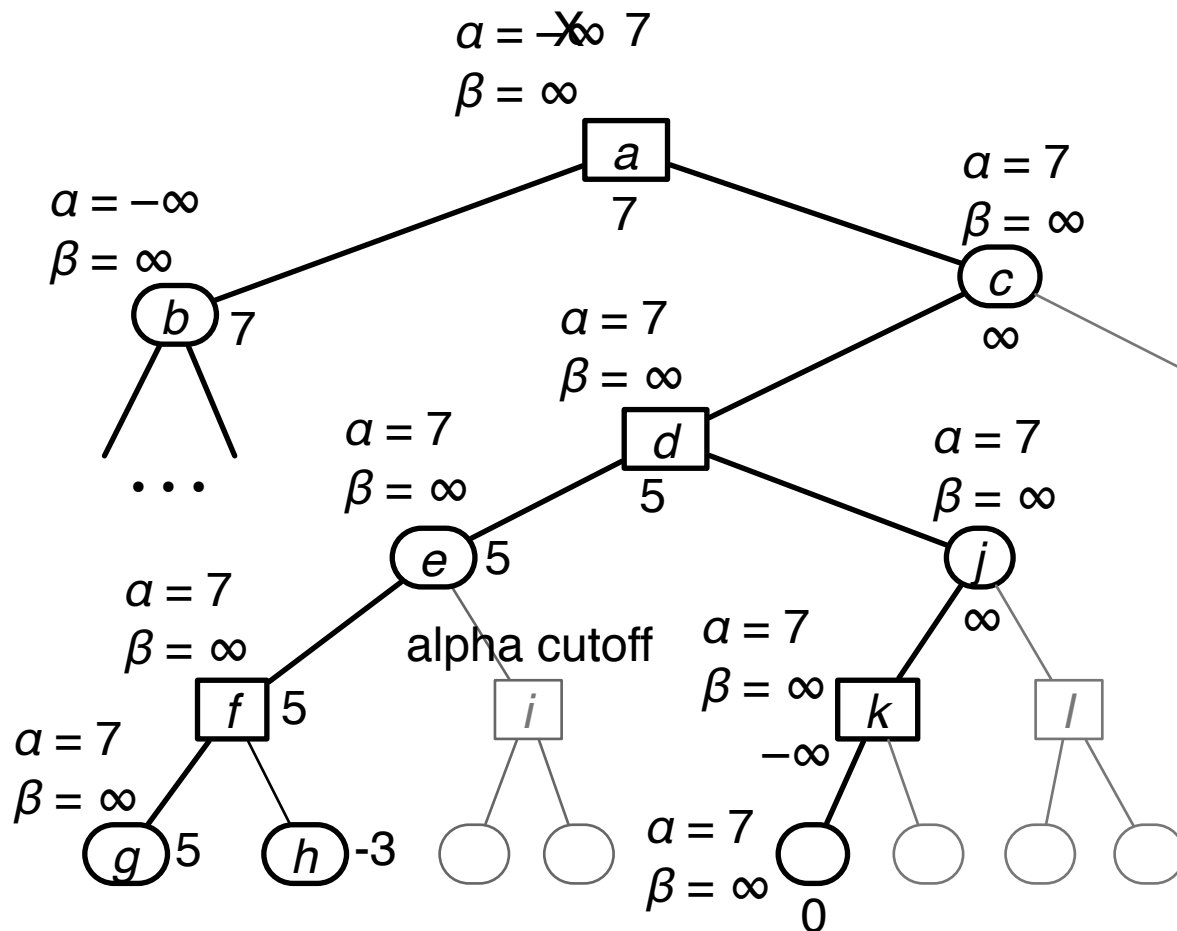
for every  $a \in \chi(h)$  do

$v \leftarrow \min(v, \text{Alpha-Beta}(\sigma(h, a), d-1, \alpha, \beta))$

if  $v \leq \alpha$  then return  $v$  //  $\alpha$  cutoff

else  $\beta \leftarrow \min(\beta, v)$

return  $v$



function Alpha-Beta( $h, d, \alpha, \beta$ )

if  $h \in Z$  then return  $u(h)$

else if  $d = 0$  then return  $e(h)$

else if  $\rho(h) = \text{Max}$  then

$v \leftarrow -\infty$

for every  $a \in \chi(h)$  do

$v \leftarrow \max(v, \text{Alpha-Beta}(\sigma(h, a), d-1, \alpha, \beta))$

if  $v \geq \beta$  then return  $v$  //  $\beta$  cutoff

else  $\alpha \leftarrow \max(\alpha, v)$

return  $v$

else

$v \leftarrow \infty$

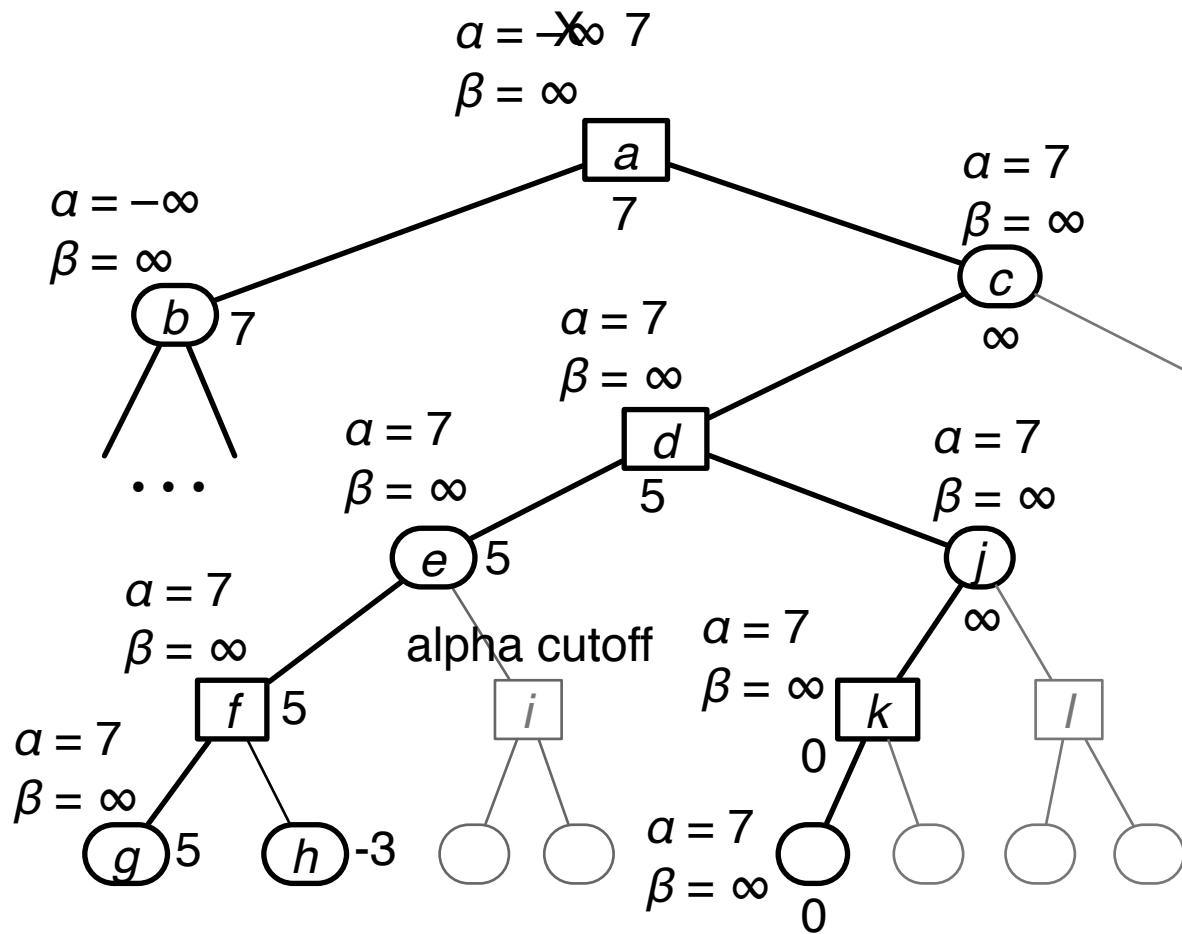
for every  $a \in \chi(h)$  do

$v \leftarrow \min(v, \text{Alpha-Beta}(\sigma(h, a), d-1, \alpha, \beta))$

if  $v \leq \alpha$  then return  $v$  //  $\alpha$  cutoff

else  $\beta \leftarrow \min(\beta, v)$

return  $v$



function Alpha-Beta( $h, d, \alpha, \beta$ )

if  $h \in Z$  then return  $u(h)$

else if  $d = 0$  then return  $e(h)$

else if  $\rho(h) = \text{Max}$  then

$v \leftarrow -\infty$

for every  $a \in \chi(h)$  do

$v \leftarrow \max(v, \text{Alpha-Beta}(\sigma(h, a), d-1, \alpha, \beta))$

if  $v \geq \beta$  then return  $v$  //  $\beta$  cutoff

else  $\alpha \leftarrow \max(\alpha, v)$

return  $v$

else

$v \leftarrow \infty$

for every  $a \in \chi(h)$  do

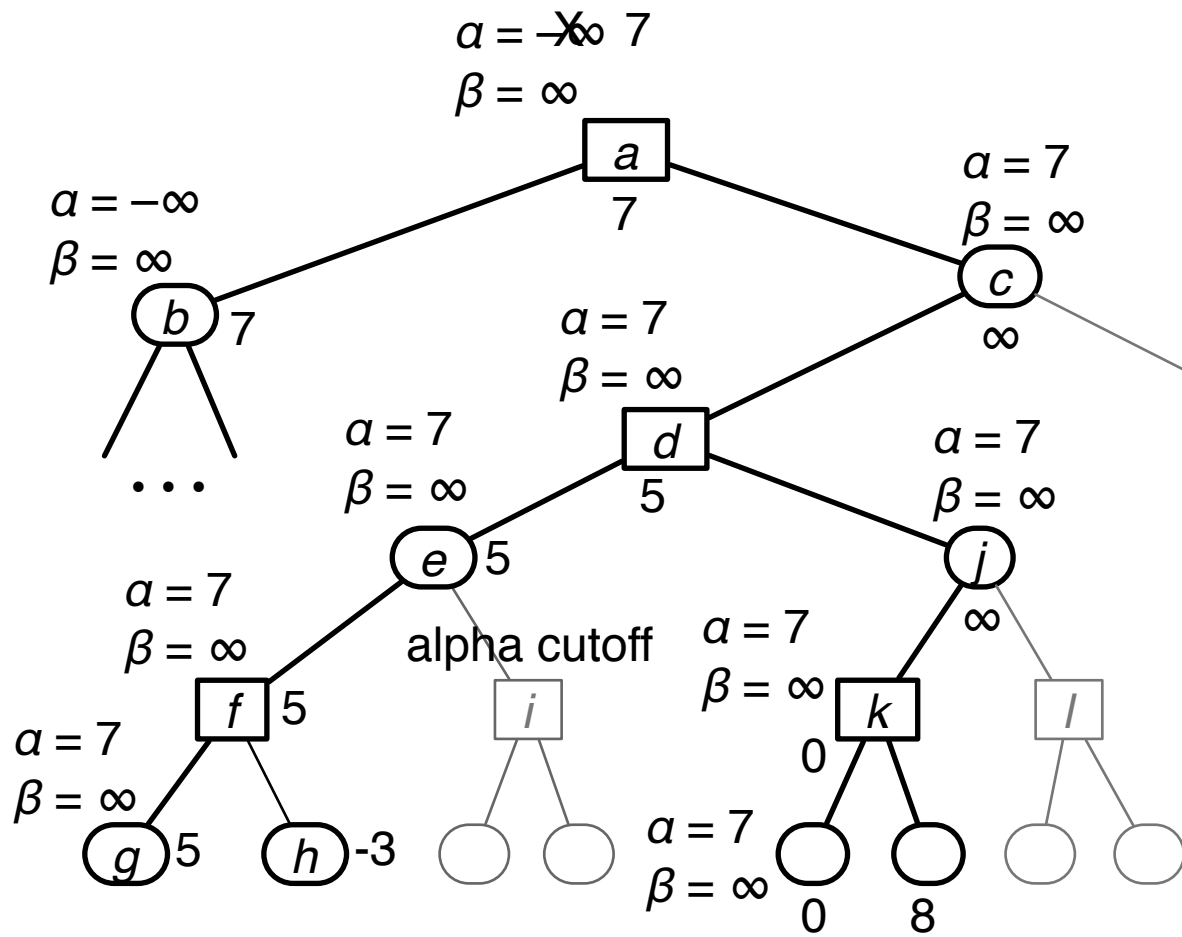
$v \leftarrow \min(v, \text{Alpha-Beta}(\sigma(h, a), d-1, \alpha, \beta))$

if  $v \leq \alpha$  then return  $v$  //  $\alpha$  cutoff

else  $\beta \leftarrow \min(\beta, v)$

return  $v$





function Alpha-Beta( $h, d, \alpha, \beta$ )

if  $h \in Z$  then return  $u(h)$

else if  $d = 0$  then return  $e(h)$

else if  $\rho(h) = \text{Max}$  then

$v \leftarrow -\infty$

for every  $a \in \chi(h)$  do

$v \leftarrow \max(v, \text{Alpha-Beta}(\sigma(h, a), d-1, \alpha, \beta))$

if  $v \geq \beta$  then return  $v$  //  $\beta$  cutoff

else  $\alpha \leftarrow \max(\alpha, v)$

return  $v$

else

$v \leftarrow \infty$

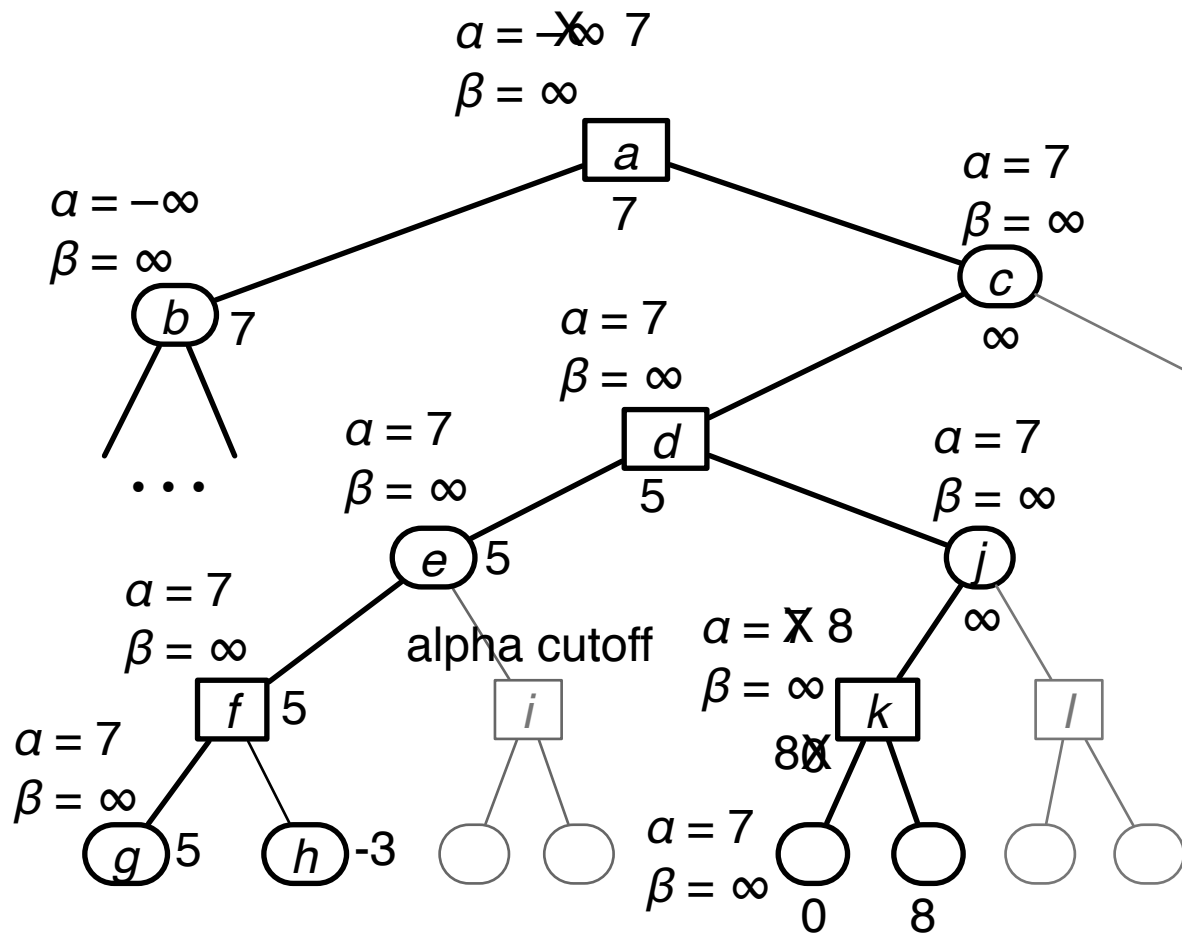
for every  $a \in \chi(h)$  do

$v \leftarrow \min(v, \text{Alpha-Beta}(\sigma(h, a), d-1, \alpha, \beta))$

if  $v \leq \alpha$  then return  $v$  //  $\alpha$  cutoff

else  $\beta \leftarrow \min(\beta, v)$

return  $v$



function Alpha-Beta( $h, d, \alpha, \beta$ )

if  $h \in Z$  then return  $u(h)$

else if  $d = 0$  then return  $e(h)$

else if  $\rho(h) = \text{Max}$  then

$v \leftarrow -\infty$

for every  $a \in \chi(h)$  do

$v \leftarrow \max(v, \text{Alpha-Beta}(\sigma(h, a), d-1, \alpha, \beta))$

if  $v \geq \beta$  then return  $v$  //  $\beta$  cutoff

else  $\alpha \leftarrow \max(\alpha, v)$

return  $v$

else

$v \leftarrow \infty$

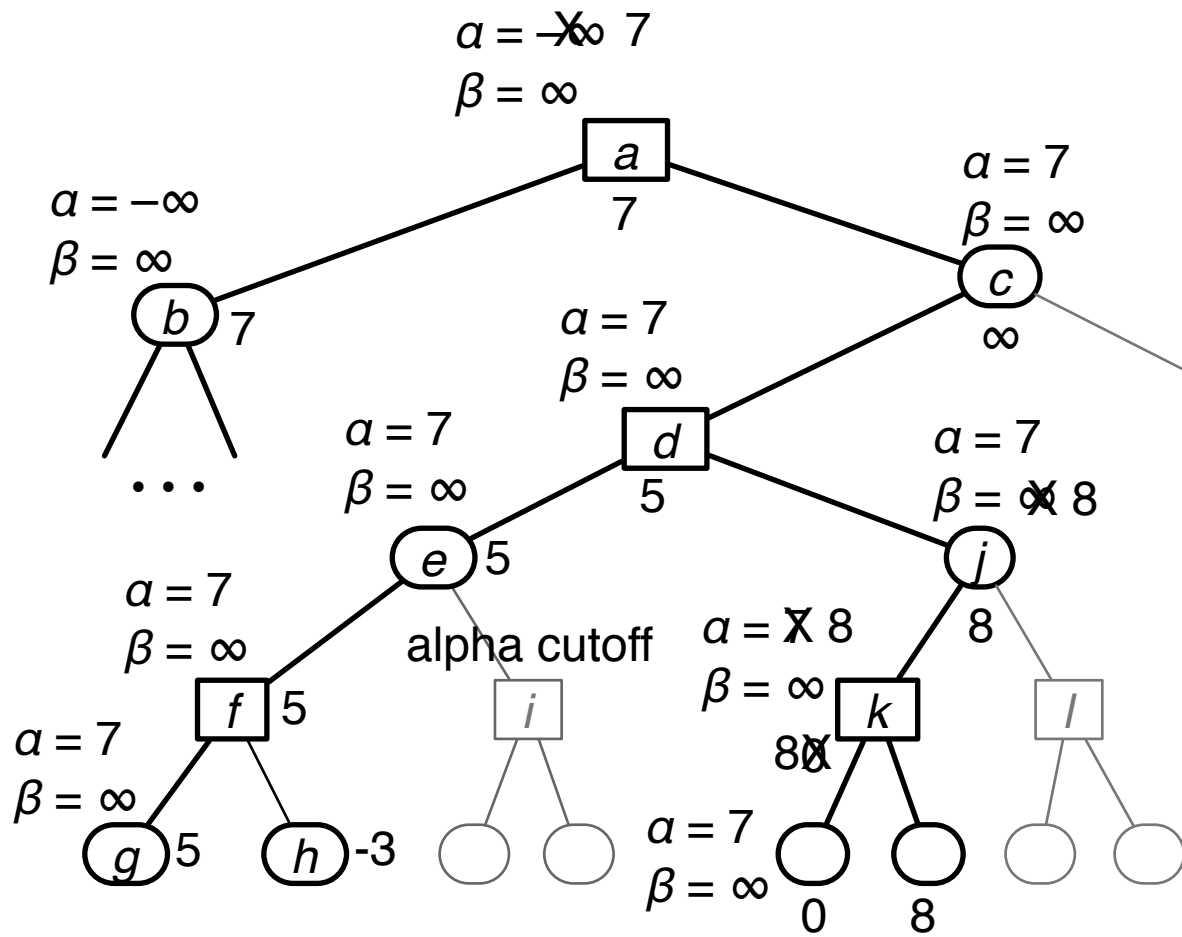
for every  $a \in \chi(h)$  do

$v \leftarrow \min(v, \text{Alpha-Beta}(\sigma(h, a), d-1, \alpha, \beta))$

if  $v \leq \alpha$  then return  $v$  //  $\alpha$  cutoff

else  $\beta \leftarrow \min(\beta, v)$

return  $v$



function Alpha-Beta( $h, d, \alpha, \beta$ )

if  $h \in Z$  then return  $u(h)$

else if  $d = 0$  then return  $e(h)$

else if  $\rho(h) = \text{Max}$  then

$v \leftarrow -\infty$

for every  $a \in \chi(h)$  do

$v \leftarrow \max(v, \text{Alpha-Beta}(\sigma(h, a), d-1, \alpha, \beta))$

if  $v \geq \beta$  then return  $v$  //  $\beta$  cutoff

else  $\alpha \leftarrow \max(\alpha, v)$

return  $v$

else

$v \leftarrow \infty$

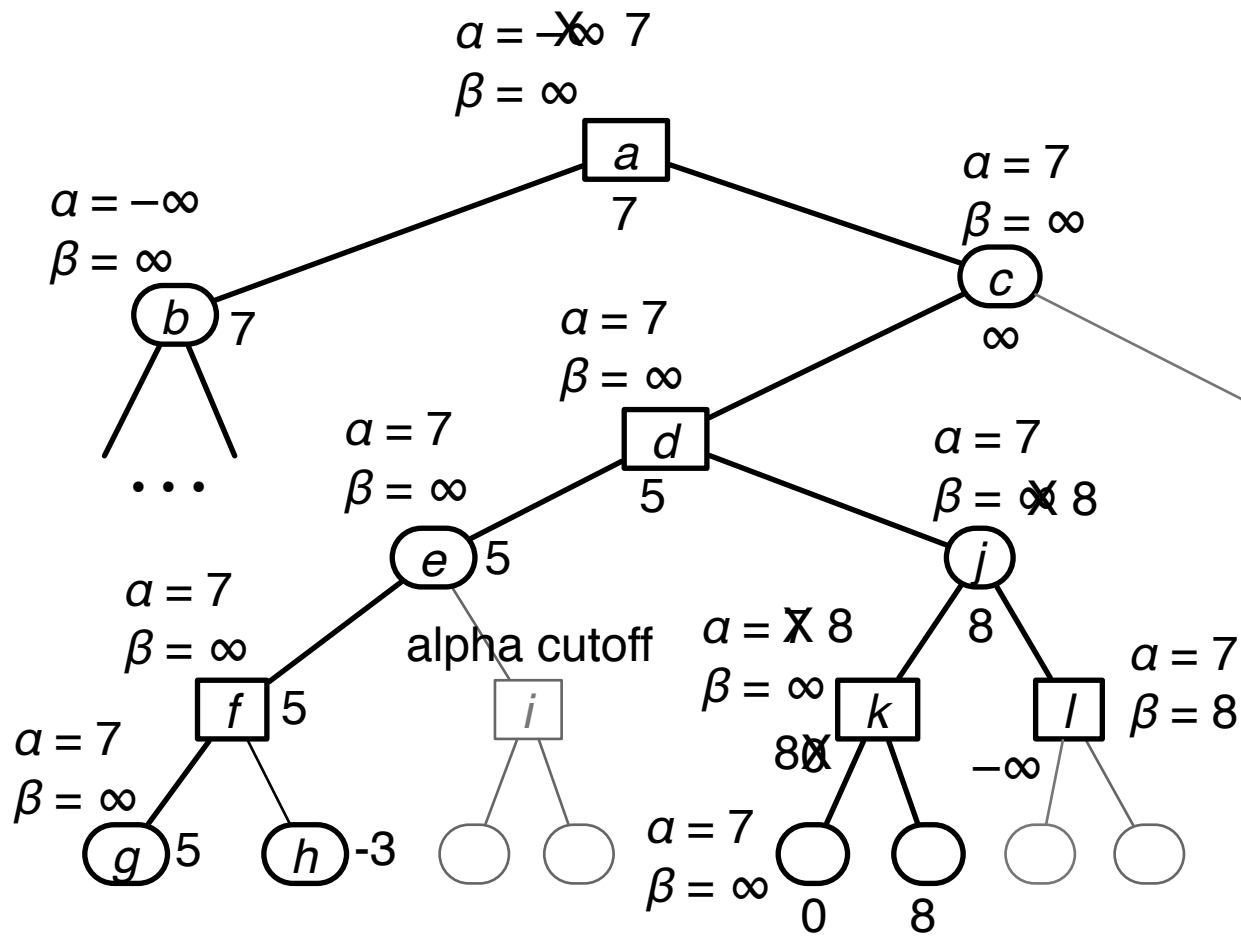
for every  $a \in \chi(h)$  do

$v \leftarrow \min(v, \text{Alpha-Beta}(\sigma(h, a), d-1, \alpha, \beta))$

if  $v \leq \alpha$  then return  $v$  //  $\alpha$  cutoff

else  $\beta \leftarrow \min(\beta, v)$

return  $v$



function Alpha-Beta( $h, d, \alpha, \beta$ )

if  $h \in Z$  then return  $u(h)$

else if  $d = 0$  then return  $e(h)$

else if  $\rho(h) = \text{Max}$  then

$v \leftarrow -\infty$

for every  $a \in \chi(h)$  do

$v \leftarrow \max(v, \text{Alpha-Beta}(\sigma(h, a), d-1, \alpha, \beta))$

if  $v \geq \beta$  then return  $v$  //  $\beta$  cutoff

else  $\alpha \leftarrow \max(\alpha, v)$

return  $v$

else

$v \leftarrow \infty$

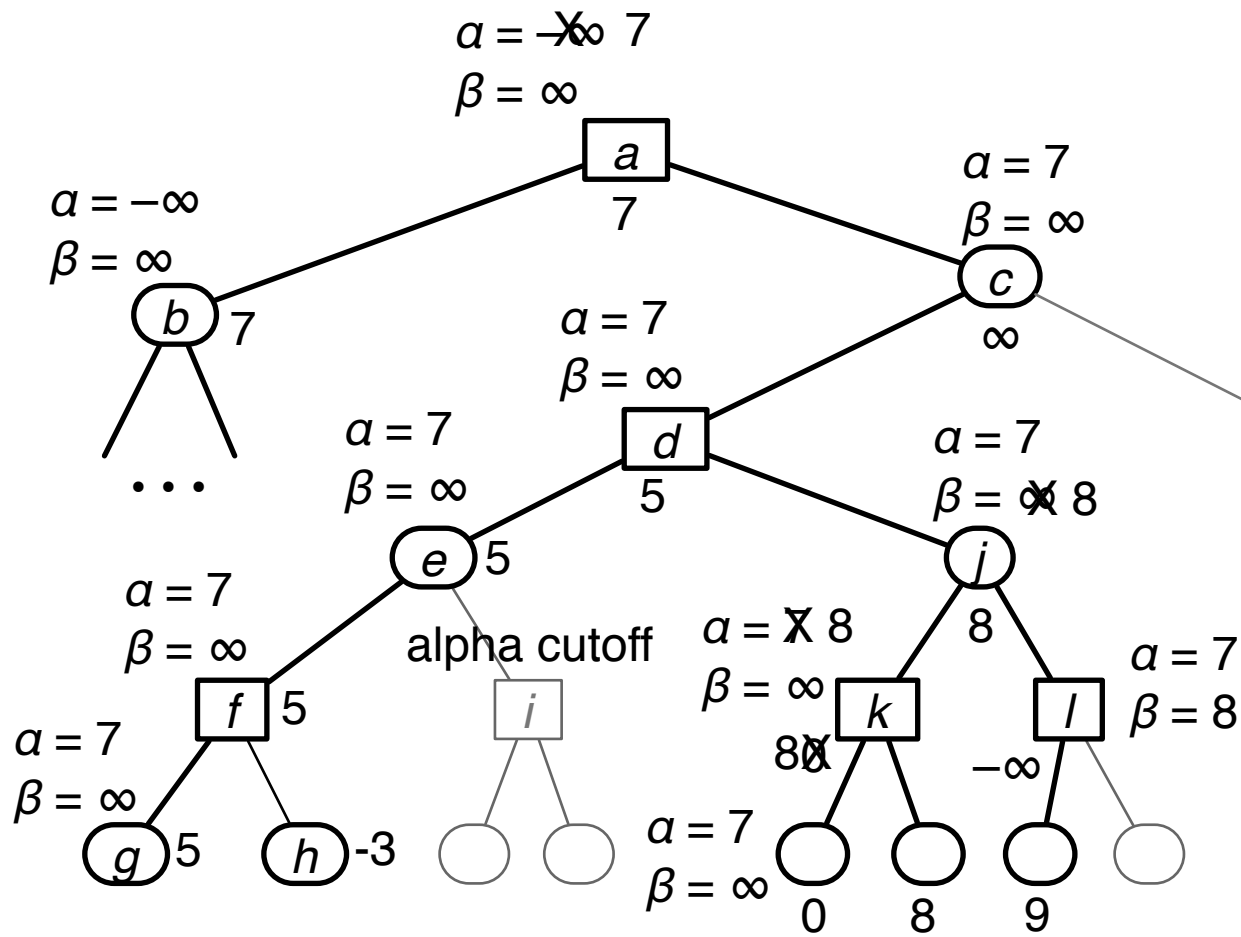
for every  $a \in \chi(h)$  do

$v \leftarrow \min(v, \text{Alpha-Beta}(\sigma(h, a), d-1, \alpha, \beta))$

if  $v \leq \alpha$  then return  $v$  //  $\alpha$  cutoff

else  $\beta \leftarrow \min(\beta, v)$

return  $v$



function Alpha-Beta( $h, d, \alpha, \beta$ )

if  $h \in Z$  then return  $u(h)$

else if  $d = 0$  then return  $e(h)$

else if  $\rho(h) = \text{Max}$  then

$v \leftarrow -\infty$

for every  $a \in \chi(h)$  do

$v \leftarrow \max(v, \text{Alpha-Beta}(\sigma(h, a), d-1, \alpha, \beta))$

if  $v \geq \beta$  then return  $v$  //  $\beta$  cutoff

else  $\alpha \leftarrow \max(\alpha, v)$

return  $v$

else

$v \leftarrow \infty$

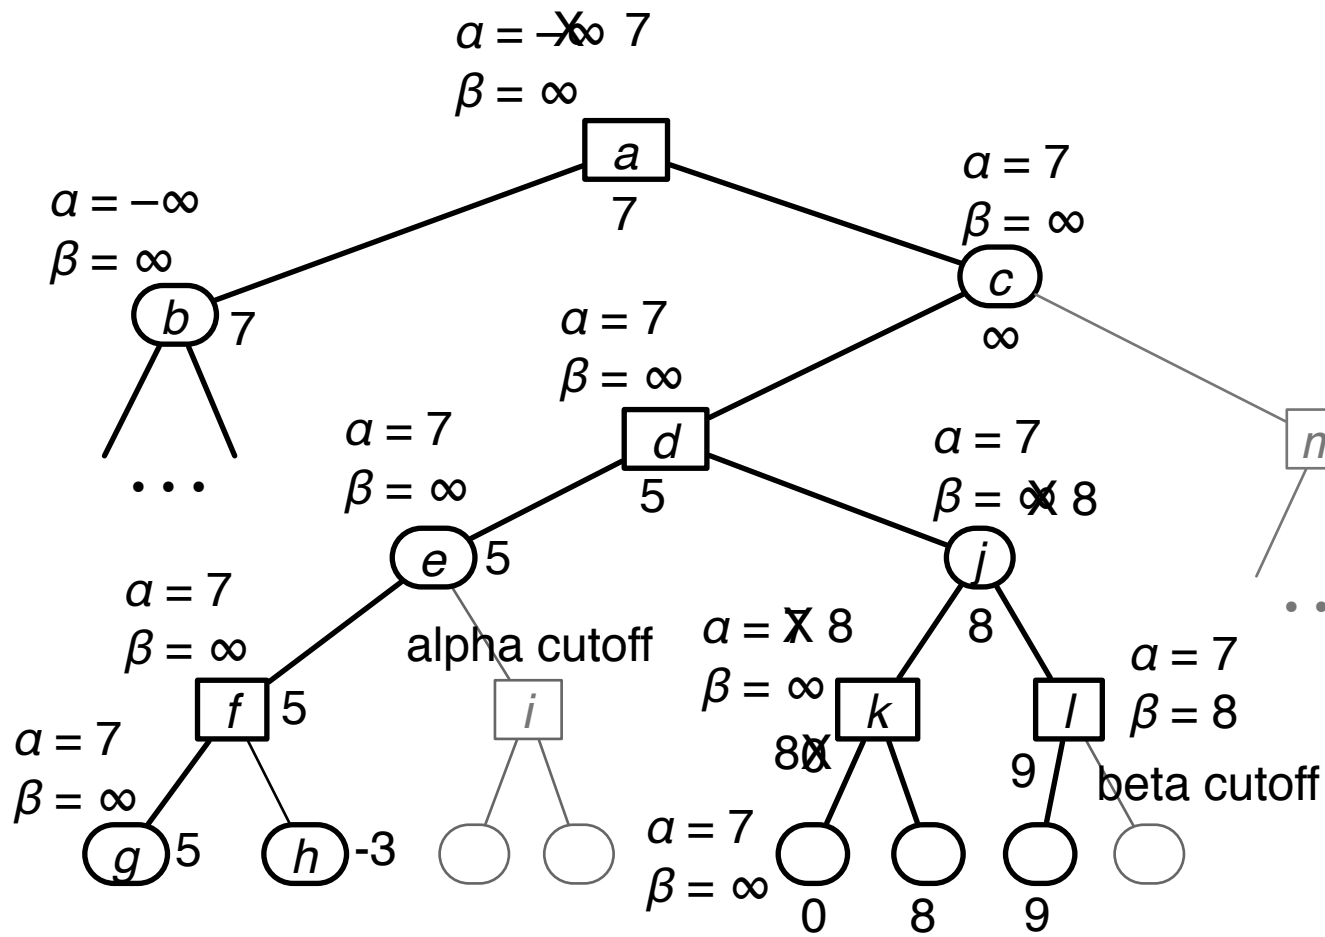
for every  $a \in \chi(h)$  do

$v \leftarrow \min(v, \text{Alpha-Beta}(\sigma(h, a), d-1, \alpha, \beta))$

if  $v \leq \alpha$  then return  $v$  //  $\alpha$  cutoff

else  $\beta \leftarrow \min(\beta, v)$

return  $v$



function Alpha-Beta( $h, d, \alpha, \beta$ )

if  $h \in Z$  then return  $u(h)$

else if  $d = 0$  then return  $e(h)$

else if  $\rho(h) = \text{Max}$  then

$v \leftarrow -\infty$

for every  $a \in \chi(h)$  do

$v \leftarrow \max(v, \text{Alpha-Beta}(\sigma(h, a), d-1, \alpha, \beta))$

if  $v \geq \beta$  then return  $v$  //  $\beta$  cutoff

else  $\alpha \leftarrow \max(\alpha, v)$

return  $v$

else

$v \leftarrow \infty$

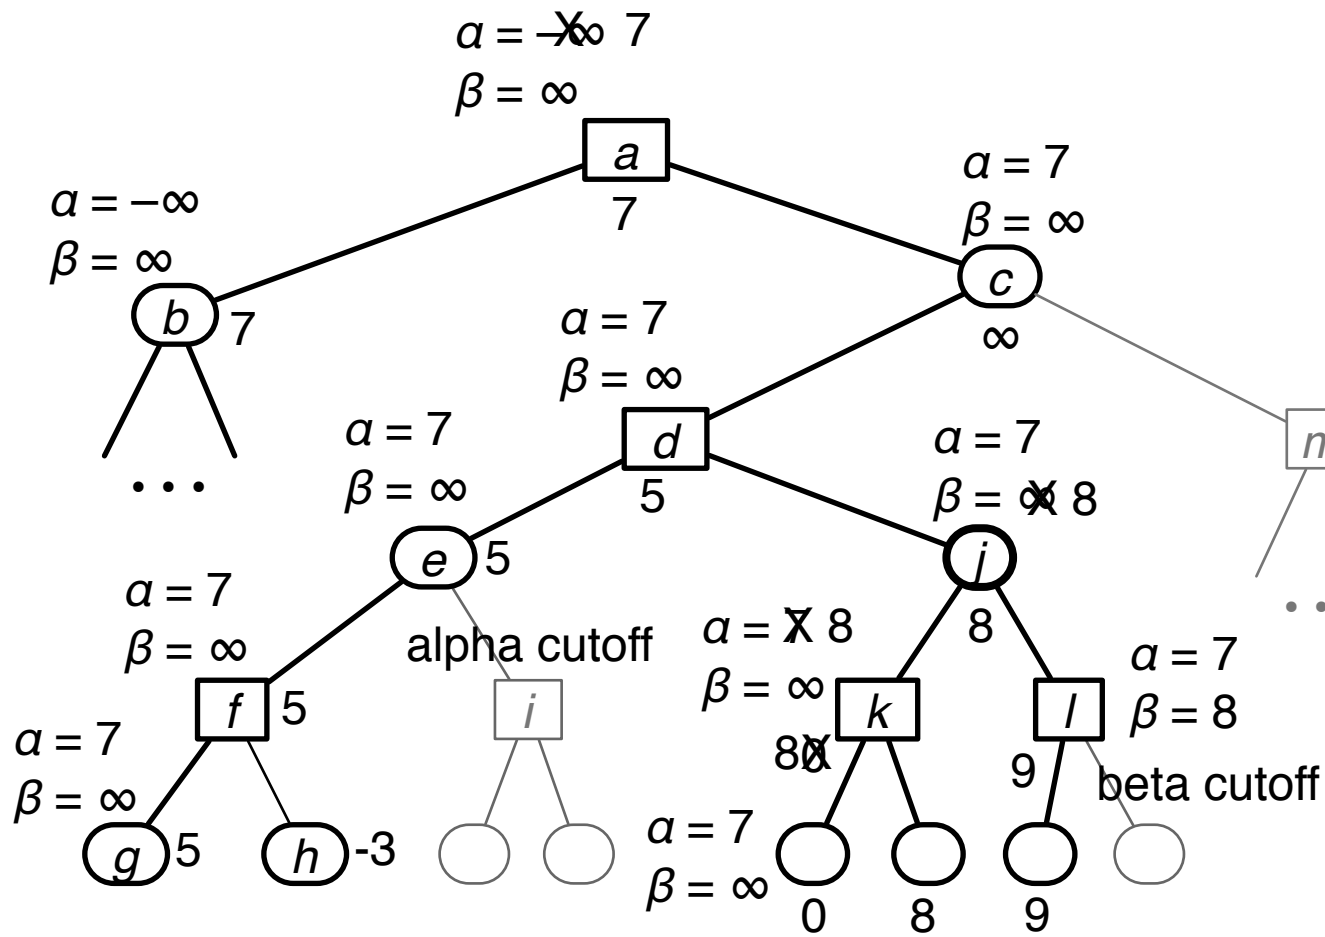
for every  $a \in \chi(h)$  do

$v \leftarrow \min(v, \text{Alpha-Beta}(\sigma(h, a), d-1, \alpha, \beta))$

if  $v \leq \alpha$  then return  $v$  //  $\alpha$  cutoff

else  $\beta \leftarrow \min(\beta, v)$

return  $v$



function Alpha-Beta( $h, d, \alpha, \beta$ )

if  $h \in Z$  then return  $u(h)$

else if  $d = 0$  then return  $e(h)$

else if  $\rho(h) = \text{Max}$  then

$v \leftarrow -\infty$

for every  $a \in \chi(h)$  do

$v \leftarrow \max(v, \text{Alpha-Beta}(\sigma(h, a), d-1, \alpha, \beta))$

if  $v \geq \beta$  then return  $v$  //  $\beta$  cutoff

else  $\alpha \leftarrow \max(\alpha, v)$

return  $v$

else

$v \leftarrow \infty$

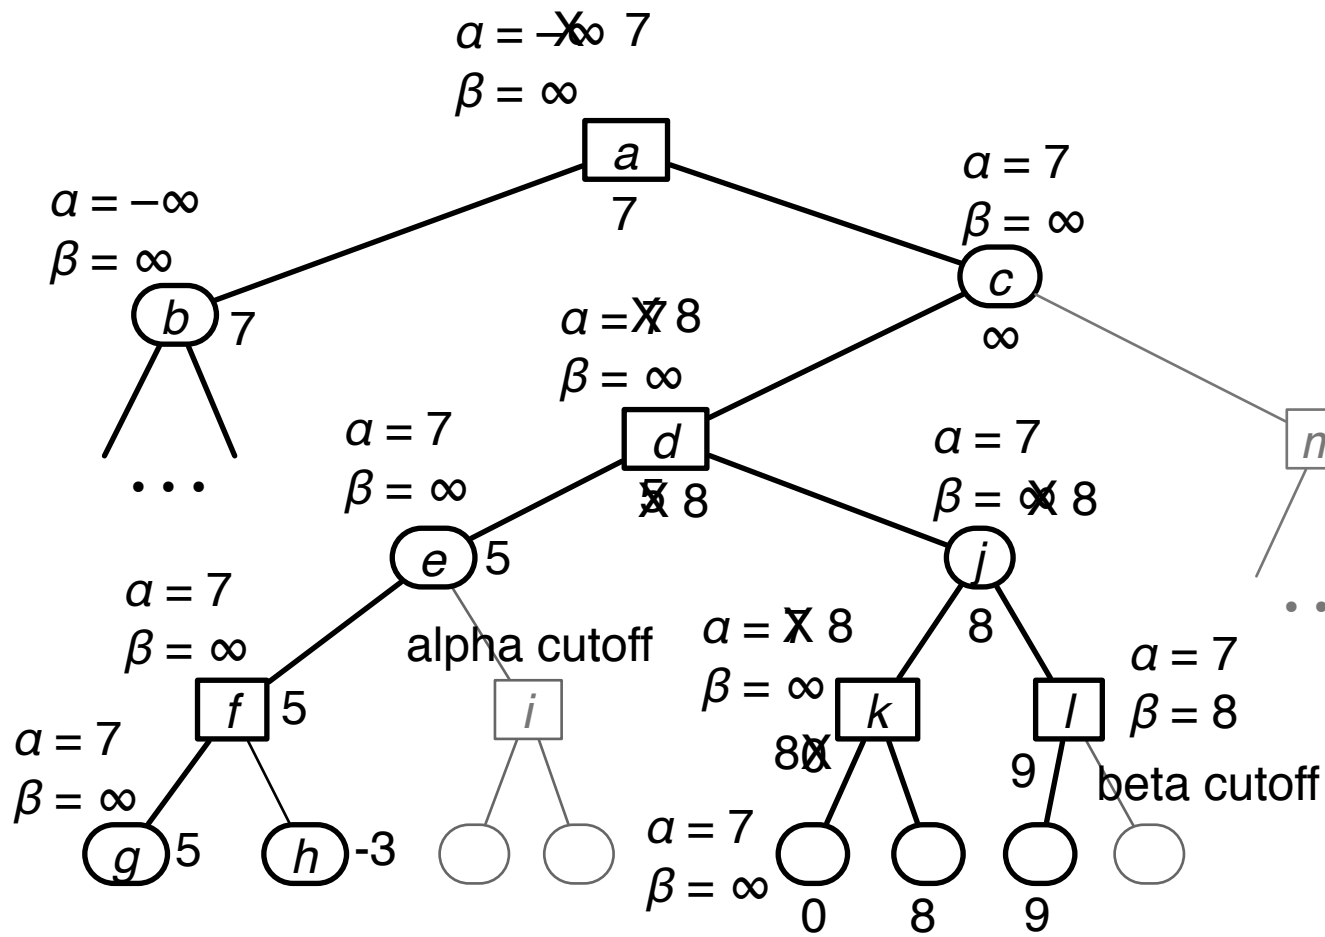
for every  $a \in \chi(h)$  do

$v \leftarrow \min(v, \text{Alpha-Beta}(\sigma(h, a), d-1, \alpha, \beta))$

if  $v \leq \alpha$  then return  $v$  //  $\alpha$  cutoff

else  $\beta \leftarrow \min(\beta, v)$

return  $v$



function Alpha-Beta( $h, d, \alpha, \beta$ )

if  $h \in Z$  then return  $u(h)$

else if  $d = 0$  then return  $e(h)$

else if  $\rho(h) = \text{Max}$  then

$v \leftarrow -\infty$

for every  $a \in \chi(h)$  do

$v \leftarrow \max(v, \text{Alpha-Beta}(\sigma(h, a), d-1, \alpha, \beta))$

if  $v \geq \beta$  then return  $v$  //  $\beta$  cutoff

else  $\alpha \leftarrow \max(\alpha, v)$

return  $v$

else

$v \leftarrow \infty$

for every  $a \in \chi(h)$  do

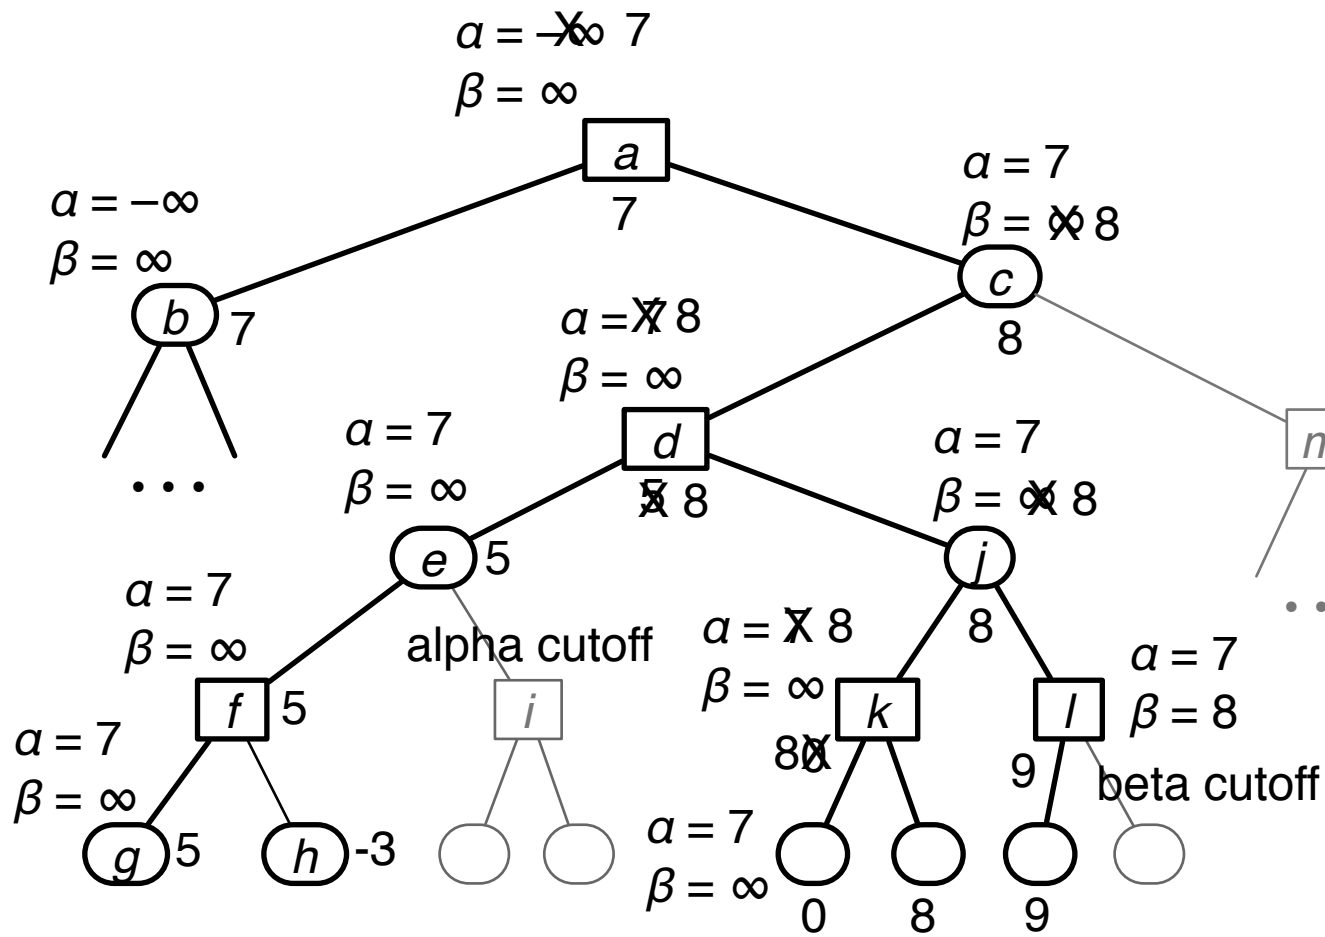
$v \leftarrow \min(v, \text{Alpha-Beta}(\sigma(h, a), d-1, \alpha, \beta))$

if  $v \leq \alpha$  then return  $v$  //  $\alpha$  cutoff

else  $\beta \leftarrow \min(\beta, v)$

return  $v$





function Alpha-Beta( $h, d, \alpha, \beta$ )

if  $h \in Z$  then return  $u(h)$

else if  $d = 0$  then return  $e(h)$

else if  $\rho(h) = \text{Max}$  then

$v \leftarrow -\infty$

for every  $a \in \chi(h)$  do

$v \leftarrow \max(v, \text{Alpha-Beta}(\sigma(h, a), d-1, \alpha, \beta))$

if  $v \geq \beta$  then return  $v$  //  $\beta$  cutoff

else  $\alpha \leftarrow \max(\alpha, v)$

return  $v$

else

$v \leftarrow \infty$

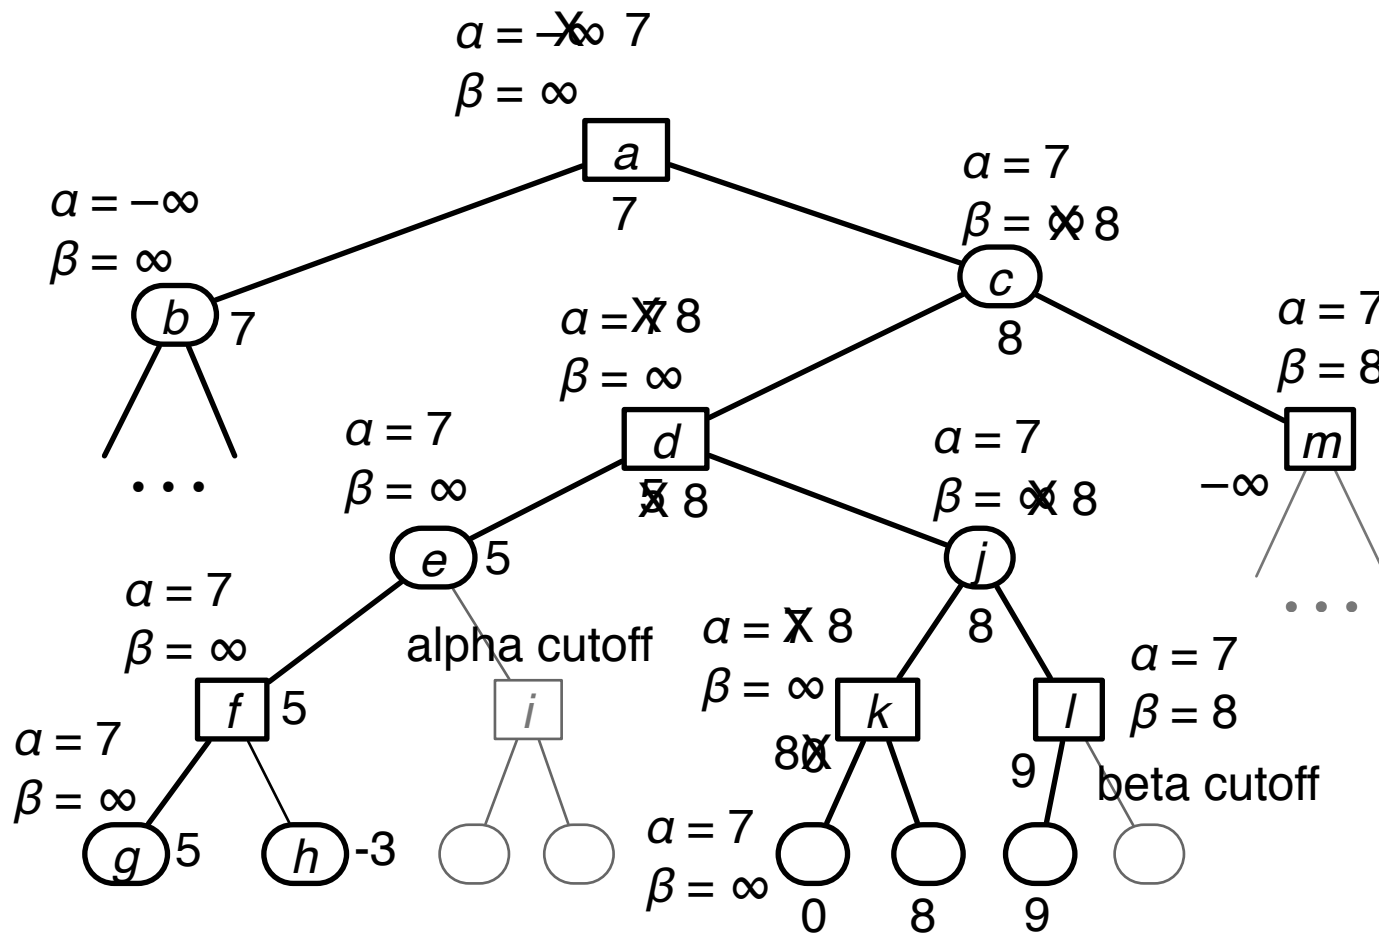
for every  $a \in \chi(h)$  do

$v \leftarrow \min(v, \text{Alpha-Beta}(\sigma(h, a), d-1, \alpha, \beta))$

if  $v \leq \alpha$  then return  $v$  //  $\alpha$  cutoff

else  $\beta \leftarrow \min(\beta, v)$

return  $v$



function Alpha-Beta( $h, d, \alpha, \beta$ )

if  $h \in Z$  then return  $u(h)$

else if  $d = 0$  then return  $e(h)$

else if  $\rho(h) = \text{Max}$  then

$v \leftarrow -\infty$

for every  $a \in \chi(h)$  do

$v \leftarrow \max(v, \text{Alpha-Beta}(\sigma(h, a), d-1, \alpha, \beta))$

if  $v \geq \beta$  then return  $v$  //  $\beta$  cutoff

else  $\alpha \leftarrow \max(\alpha, v)$

return  $v$

else

$v \leftarrow \infty$

for every  $a \in \chi(h)$  do

$v \leftarrow \min(v, \text{Alpha-Beta}(\sigma(h, a), d-1, \alpha, \beta))$

if  $v \leq \alpha$  then return  $v$  //  $\alpha$  cutoff

else  $\beta \leftarrow \min(\beta, v)$

return  $v$

# Properties of Alpha-Beta

- Alpha-beta pruning reasons about which computations are relevant
  - A form of **metareasoning**

## Theorem:

- If  $\text{LD-minimax}(h, d)$  returns a value in  $[\alpha, \beta]$ 
  - $\text{Alpha-Beta}(h, d, \alpha, \beta)$  returns the same value
- If  $\text{LD-minimax}(h, d)$  returns a value  $\leq \alpha$ 
  - $\text{Alpha-Beta}(h, d, \alpha, \beta)$  returns a value  $\leq \alpha$
- If  $\text{LD-minimax}(h, d)$  returns a value  $\geq \beta$ 
  - $\text{Alpha-Beta}(h, d, \alpha, \beta)$  returns a value  $\geq \beta$

## Corollary:

- $\text{Alpha-Beta}(h, d, -\infty, \infty)$  returns the same value as  $\text{LD-minimax}(h, d)$
- $\text{Alpha-Beta}(h, \infty, -\infty, \infty)$  returns  $u(h)$

# Node Ordering

- Deeper search (larger  $d$ ) usually gives better decisions

- There are “pathological” games where it doesn’t, but those are rare


- How much deeper can Alpha-Beta search?

- Worst case:

- children of Max nodes are smallest-value-first
  - children of Min nodes are greatest-value-first
  - Alpha-Beta visits all nodes of depth  $\leq d$ 
    - running time  $O(b^d)$

- Best case:

- children of Max nodes are greatest-value-first
  - children of Min nodes are smallest-value-first
  - Alpha-Beta’s running time is  $O(b^{d/2}) \Rightarrow$  doubles the solvable depth



# Node Ordering

- How to get closer to the best case:
  - Every time you expand a state  $s$ , apply  $e$  to its children
  - When it's Max's move, sort the children in order of largest  $e$  first
  - When it's Min's move, sort the children in order of smallest  $e$  first
- Suppose we have 100 seconds, explore  $10^4$  nodes/second
  - $10^6$  nodes per move
  - Chess midgame  $b \approx 35$ 
    - LD-minimax: time  $b^d = 35^d = 10^6 \rightarrow d \approx 4$
    - Alpha-Beta: time  $b^{d/2} \approx 35^{8/2} \rightarrow d \approx 8$

# Other Modifications

- Several other ways to improve accuracy or running time:
  - quiescence search and biasing
  - transposition tables
  - thinking on the opponent's time
  - table lookup of “book moves”
  - iterative deepening
  - forward pruning

# Quiescence Search and Biasing

- In a game like checkers or chess,  $e$  is based greatly on material pieces
  - Likely to be inaccurate if there are pending captures
- Search deeper to reach a position where there aren't pending captures
  - $e$  will be more accurate here
- That creates another problem
  - You're searching some paths to an even depth, others to an odd depth
  - Paths that end just after your opponent's move will look worse than paths that end just after your move
- To compensate, add or subtract a number called the “biasing factor”


# Transposition Tables

- Multiple paths to the same state (state space is a graph rather than a tree)
- Idea:
  - when you compute  $s$ 's minimax value, store it in a hash table
  - visit  $s$  again  $\Rightarrow$  retrieve its value rather than computing it again
- The hash table is called a **transposition table**
- Problem: can't store exponentially many states
  - Only store some – the ones that you're most likely to need



# Thinking on the Opponent's Time

- Suppose you're at node  $a$ 
  - Use Alpha-Beta to estimate  $u(b)$  and  $u(c)$
  - $c$  looks better, so move there



- Consider your estimates of  $u(f)$  and  $u(g)$ 
  - They suggest your opponent is likely to move to  $f$
  - While waiting for the opponent to move, start an alpha-beta search below  $f$
- If the opponent moves to  $f$ , then you've already done a lot of the work of figuring out your next move

# Book Moves

- In some games, experts have spent lots of time analyzing **openings**
  - Sequences of moves one might play at the start of the game
  - Best responses to those sequences
- Some of these are cataloged in standard reference works
  - e.g., the *Encyclopaedia of Chess Openings*
- Store these in a lookup table
  - Respond almost immediately, as long as the opponent sticks to a sequence that's in the book
- A technique humans can use when playing against such a program
  - Deliberately make a move that isn't in the book
  - This may weaken the human's position, but the computer will (i) start taking longer and (ii) stop playing as well

# Iterative Deepening

- How deeply should you search a game tree?
  - When you call  $\text{Alpha-Beta}(h, d, -\infty, \infty)$ , what to use for  $d$ ?
- small  $d \Rightarrow$  don't make as good a decision
- large  $d \Rightarrow$  run out of time without knowing what move to make
- Solution: **iterative deepening**

for  $d = 1$  by 1 until you run out of time


$m \leftarrow$  the move returned by  $\text{Alpha-Beta}(h, d, -\infty, \infty)$

- Why this works:
  - Time complexity is  $O(b^1 + b^2 + \dots + b^d) = O(b^d)$
  - For large  $b$ , each iteration takes much more time than the total of *all the previous iterations*

# Forward Pruning

- **Tapered search:**

- Instead of looking at all of a node's children, just look at the  $n$  best ones
  - the  $n$  highest  $e(h)$  values
- Decrease the value of  $n$  as you go deeper in the tree
- Drawback: may exclude an important move at a low level of the tree



- **Marginal forward pruning:**

- Ignore every child  $h$  such that  $e(h)$  is smaller than the current best values from the nodes you've already visited
- Not reliable, should be avoided

# Forward Pruning

- Until the mid-1970s most computer-chess researchers tried to get programs to search “like people think”
  - Use extensive chess knowledge at each node to select a few “plausible” moves; prune the others
  - Serious tactical shortcomings
- Brute-force search programs did better, dominated computer chess for about 20 years
- Early 1990s: development of some forward-pruning techniques that worked well
  - Null-move pruning
- Today most chess programs use some kind of forward-pruning
  - null-move pruning is one of the most popular

# Game-Tree Search in Practice

- **Othello:** since 1980, the best computer programs have easily defeated the best humans
- **Checkers:** In 1994, Chinook ended 40-year-reign of human world champion Marion Tinsley
  - Tinsley withdrew for health reasons, died a few months later
- In 2007, Checkers was solved
  - With perfect play, it's a draw
  - This took  $10^{14}$  calculations over 18 years
  - Search space size  $5 \times 10^{20}$
- **Chess:** In 1997, Deep Blue defeated Gary Kasparov in a six-game match
  - Used special hardware, could search 200 million positions per second
  - Modern programs don't use special hardware
  - Barred from world-championship matches (except against other computers)

# Game-Tree Search in Practice

- **Go:** before 2006, good amateurs could easily beat the best go programs, even when given a handicap
- Two *big* jumps in performance
  - 2006: Monte Carlo rollouts
    - brought go programs up to a master level, better than strong amateurs
  - 2015: Monte Carlo rollouts combined with deep learning
    - AlphaGo: first go program to beat professional human players

# Rules of Go (Abbreviated)


- Go-board:  $19 \times 19$  locations (intersections on a grid)
- Black and White take turns
- Black has the first move
- Each move consists of placing a stone at an unoccupied location
  - You just put them on the board; you don't move them around
- Adjacent stones of the same color are called a **string**.
  - **Liberties** are the empty locations next to the string
  - A string is removed if its number of liberties is 0
- Score: territories (number of occupied or surrounded locations)





# Why Go is Difficult for Computers

- A game tree's size grows exponentially with both its depth and its branching factor
- The game tree for go:
  - branching factor  $\approx 200$
  - game length  $\approx 250$  to  $300$  moves
  - number of nodes in the game tree  $\approx 10^{525}$  to  $10^{620}$ 
    - vs.  $10^{135}$  for chess
- Much too big for a normal game tree search



# Why Go is Difficult for Computers


- Writing an evaluation function for chess
  - Mainly piece count, plus some positional considerations
    - isolated/doubled pawns, rooks on open files (columns), pawns in the center of the board, etc.
  - As the game progresses, pieces get removed => evaluation gets easier
- For Go, much more complicated
  - whether or not a group is alive
  - which stones can be connected to one another
  - the extent to which a position has influence or can be attacked
  - As the game progresses, pieces get added => evaluation gets more complicated

# Monte Carlo Roll-Outs

- Basic idea:
  - Generate a list of potential moves  $A = \{a_1, \dots, a_n\}$
  - For each move  $a_i$  in  $A$ :
    - Starting at  $\sigma(h, a_i)$ , generate a set of random games  $G_i$  in which the two players make all their moves at random
  - Choose the move in  $A$  that produces the best results
- Whether this works depends on **how** you generate the random games


# Multi-Arm Bandit

- Statistical model of sequential experiments
  - Name comes from a traditional slot machine (one-armed bandit)
- Multiple actions
  - Each action provides a reward from a probability distribution associated with that specific action
  - Objective: maximize the expected utility of a sequence of actions
- Exploitation vs exploration dilemma:
  - **Exploitation**: choosing an action that you already know about, because you think it's likely to give you a high reward
  - **Exploration**: choosing an action that you don't know much about, in hopes that maybe it will produce a better reward than the actions you already know about



# UCB (Upper Confidence Bound) Algorithm

- Let
  - $r_i$  = average reward you've gotten from arm  $i$
  - $t_i$  = number of times you've tried arm  $i$ ;
  - $t$  = total number of tries =  $\sum_i t_i$
- **loop**
  - **if** there are one or more arms that have not been played
  - **then** play one of them
  - **else** play the arm  $i$  that has the highest value of  $r_i + 2 \sqrt{(\log t)/t_i}$



# UCT (UCB for Trees)

- Adaptation of UCB for game-tree search
- First time it was used in go: Mogo, 2006
  - Won the go tournament at the 2007 Computer Olympiad, and several other computer go tournaments
- Now used in most computer go programs
- Next page: basic idea (omitting several improvements)

global  $Seen \leftarrow \emptyset$

## UCT (Basic Idea)

function UCT( $h$ )

if  $h \notin Seen$  then

add  $h$  to  $Seen$

$r_h \leftarrow 0$ ;  $t_h \leftarrow 0$

$h' \leftarrow \text{UCB-choose}(\{\sigma(h, a) \mid a \in \chi(h)\})$

$\mathbf{v} \leftarrow \text{UCT}(h')$

$r_h \leftarrow (r_h t_h + \mathbf{v}[\rho(h)]) / (t_h + 1)$

$t_h \leftarrow t_h + 1$

return  $\mathbf{v}$

$Seen = \{\text{all nodes seen so far}\}$

$t_h = \text{no. of times we've tried } h$

$r_h = \text{avg. reward for } \rho(h) \text{ at } h$

$\rho(h) = \text{player to move at } h$

$\chi(h) = \{\text{available actions at } h\}$

$\sigma(h, a) = \text{child of } h \text{ produced by } a$

function UCB-choose( $H$ )

if  $H \setminus Seen \neq \emptyset$  then return a randomly chosen member of  $H \setminus Seen$

else

$t \leftarrow \sum_{h \in H} t_h$

return the  $h \in H$  with the highest value of  $r_h + 2 \sqrt{(\log t) / t_h}$

global  $Seen \leftarrow \emptyset$

## UCT (Basic Idea)

function UCT( $h$ )

if  $h \notin Seen$  then

add  $h$  to  $Seen$

$r_h \leftarrow 0$ ;  $t_h \leftarrow 0$

$h' \leftarrow \text{UCB-choose}(\{\sigma(h, a) \mid a \in \chi(h)\})$

$\mathbf{v} \leftarrow \text{UCT}(h')$

$r_h \leftarrow (r_h t_h + \mathbf{v}[\rho(h)]) / (t_h + 1)$

$t_h \leftarrow t_h + 1$

return  $\mathbf{v}$

- As the number of iterations  $\rightarrow \infty$ , each  $r_h \rightarrow u(h)$

function UCB-choose( $H$ )

if  $H \setminus Seen \neq \emptyset$  then return a randomly chosen member of  $H \setminus Seen$

else

$t \leftarrow \sum_{h \in H} t_h$

return the  $h \in H$  with the highest value of  $r_h + 2 \sqrt{(\log t) / t_h}$



global  $Seen \leftarrow \emptyset$

## UCT (Basic Idea)

function UCT( $h$ )

if  $h \notin Seen$  then

add  $h$  to  $Seen$

$r_h \leftarrow 0$ ;  $t_h \leftarrow 0$

$h' \leftarrow \text{UCB-choose}(\{\sigma(h, a) \mid a \in \chi(h)\})$

$\mathbf{v} \leftarrow \text{UCT}(h')$

$r_h \leftarrow (r_h t_h + \mathbf{v}[\rho(h)]) / (t_h + 1)$

$t_h \leftarrow t_h + 1$

return  $\mathbf{v}$

- As the number of iterations  $\rightarrow \infty$ , each  $r_h \rightarrow u(h)$
- Question: how can this be true?
  - $u(h)$  should mean each player chooses their max payoff
  - But  $r_h$  is *average* payoff

function UCB-choose( $H$ )

if  $H \setminus Seen \neq \emptyset$  then return a randomly chosen member of  $H \setminus Seen$

else

$t \leftarrow \sum_{h \in H} t_h$

return the  $h \in H$  with the highest value of  $r_h + 2 \sqrt{(\log t) / t_h}$

global  $Seen \leftarrow \emptyset$

# UCT (Basic Idea)

function UCT( $h$ )

if  $h \notin Seen$  then

add  $h$  to  $Seen$

$r_h \leftarrow 0$ ;  $t_h \leftarrow 0$

$h' \leftarrow \text{UCB-choose}(\{\sigma(h, a) \mid a \in \chi(h)\})$

$\mathbf{v} \leftarrow \text{UCT}(h')$

$r_h \leftarrow (r_h t_h + \mathbf{v}[\rho(h)]) / (t_h + 1)$

$t_h \leftarrow t_h + 1$

return  $\mathbf{v}$

function UCB-choose( $H$ )

if  $H \setminus Seen \neq \emptyset$  then return a randomly chosen member of  $H \setminus Seen$

else

$t \leftarrow \sum_{h \in H} t_h$

return the  $h \in H$  with the highest value of  $r_h + 2 \sqrt{(\log t) / t_h}$


- As the number of iterations  $\rightarrow \infty$ , each  $r_h \rightarrow u(h)$
- Question: how can this be true?
  - $u(h)$  should mean each player chooses their max payoff
  - But  $r_h$  is *average* payoff
- Answer: as iterations  $\rightarrow \infty$ ,  $\sqrt{(\log t) / t_h} \rightarrow 0$ 
  - UCB-choose converges to choosing *highest* avg. payoff

# AlphaGo: Game Playing

- 2015: first go program to beat a professional go player
- 2016: first program to beat one of the world's top go players
- Uses a combination of Monte-Carlo with artificial neural networks


# Artificial Neural Networks

- Collection of “units” that act a little like biological neurons
- Each unit:




# Artificial Neural Networks

- Generally organized into layers
- Can train them to compute a wide variety of functions
  - Give them lots of examples: (*input*, *output*) pairs
- Formulas for adjusting the bias weights to try to do well on all the pairs



# Artificial Neural Networks

- *Convolutional* neural network: a particular way to organize the layers
- *Deep* neural network: many layers
- *Deep learning*: learning using a deep neural network



# AlphaGo: Training

- Policy network
  - Give it examples of moves experts would make in various situations
    - (*input, output*) pair: (position, expert move)
  - Resulting neural network:
    - input: position → output: probability distribution over moves
  - Do some optimization to make it work better
    - I'll skip the details
- Value network
  - Use the policy network to play games against itself
  - Use the results to train another deep neural network
    - (*input, output*) pair: (position, who won)
  - Resulting neural network:
    - input: position → output: expected utility


# AlphaGo: Training

- Generate Monte Carlo rollouts
  - Somewhat like UCT – but in addition to  $r_h$  and  $t_h$  it involves
    - probability of  $h$  returned by the neural network
    - value of  $h$  returned by the value network
- I don't know the details



# Multiplayer Games

- **Max<sup>n</sup>** algorithm: backward induction for  $n$  players, with cutoff depth  $d$  and evaluation function  $\mathbf{e}$
- Node evaluation is a payoff profile  $\mathbf{v}$   
 $\mathbf{v}[i]$  is player  $i$ 's payoff
- Approximate SPE payoff profile
  - Exact if  $d \geq \text{height of } h$



function  $\text{Maxn}(h, d)$

if  $h \in Z$  then return  $\mathbf{u}(h)$

if  $d = 0$  then return  $\mathbf{e}(h)$

$V = \{\text{Maxn}(\sigma(h, a), d-1) \mid a \in \chi(h)\}$

return  $\arg \max_{\mathbf{v} \in V} \mathbf{v}[\rho(h)]$

function  $\text{Maxn-choice}(h, d)$

if  $h \in Z$  or  $d = 0$  then return *error*

return  $\arg \max_{a \in \chi(h)} (\text{Maxn}(\sigma(h, a), d-1))[\rho(h)]$

$H = \{\text{nonterminal nodes}\}$

$Z = \{\text{terminal nodes}\}$

$\rho(h)$  = the player to move at  $h$

$\chi(h) = \{\text{available actions at } h\}$


$\sigma(h, a)$  = child of  $h$  produced by  $a$

$\mathbf{u}(h)$  = utility profile for  $h$

$\mathbf{v}[i]$  =  $i$ 'th element of  $\mathbf{v}$

# Multiplayer Games

- The **Paranoid** algorithm
  - Cutoff depth  $d$  and evaluation function  $e$
  - At  $i$ 's move take max, at elsewhere take min
- Approximate maxmin value for  $i$ 
  - exact value if  $d \geq \text{height of } h$



function Paranoid( $i, h, d$ )

if  $h \in Z$  then return  $u_i(h)$

if  $d = 0$  then return  $e_i(h)$

if  $\rho(h) = i$  then return  $\max_{a \in \chi(h)} \text{Paranoid}(i, \sigma(h, a), d-1)$

else return  $\min_{a \in \chi(h)} \text{Paranoid}(i, \sigma(h, a), d-1)$

function Paranoid-choice ( $i, h, d$ )

if  $\rho(h) = i$  then return  $\arg \max_{a \in \chi(h)} \text{Paranoid}(i, \sigma(h, a), d-1)$



else return *error*

# Discussion

- Neither Max<sup>n</sup> nor Paranoid has been entirely successful
- Partly due to dynamic relationships among players
  - Human players may hold grudges
  - Informal alliances form and dissolve over time
    - players who are behind may “gang up” on a player who’s ahead
- Max<sup>n</sup> and Paranoid don’t model these relationships
  - But they can greatly influence the players’ strategies
- For better play in a multi-player game, need ways
  - Decide when/whether to cooperate with others
  - Deduce each player’s attitude toward the other players

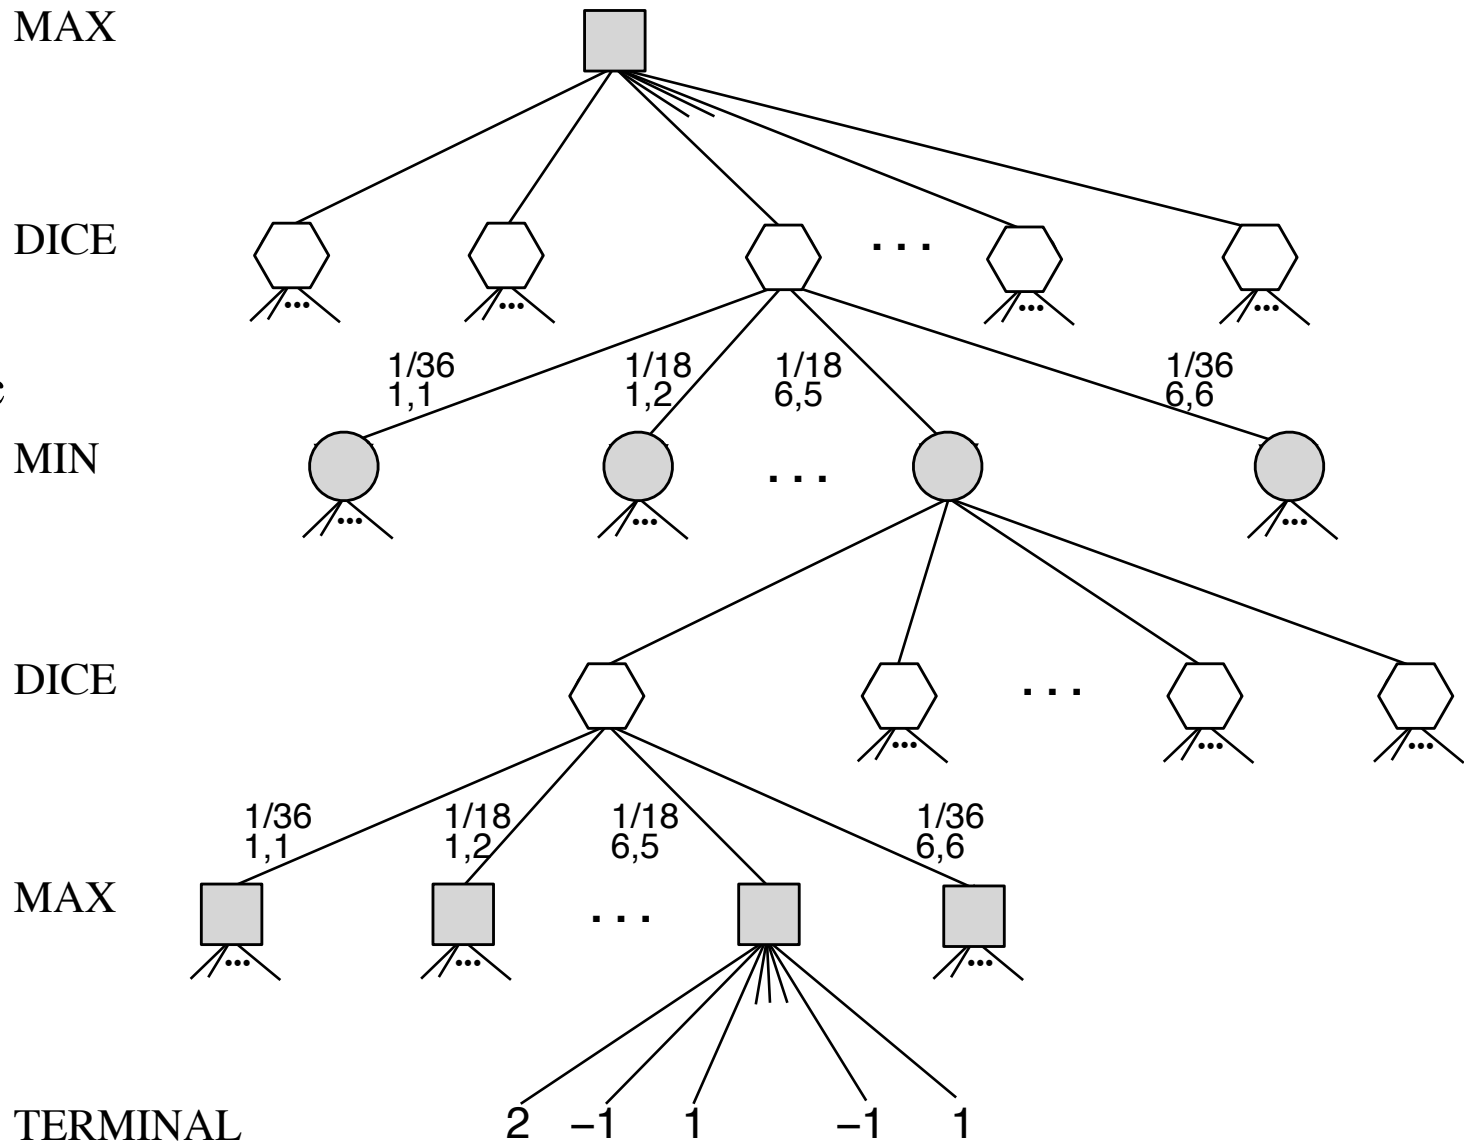
# Games with Chance Nodes

- Example: Backgammon
- Two agents who take turns
- At each turn, the set of available moves depends on the results of rolling the dice
  - Each die specifies how far to move one of your pieces
    - (except if you roll doubles)
  - You can't move to a location where your opponent has 2 or more pieces
  - You can move to a location where your opponent has 1 piece
    - Knock that piece off the board, it must start over



# Backgammon Game Tree

- Players' moves have deterministic outcomes
- Dice rolls have stochastic outcomes



# ExpectiMinimax

function ExpectiMinimax( $h, d$ )

if  $h \in Z$  then return  $u(h)$

else if  $d = 0$  then return  $e(h)$

else if  $\rho(h) = \text{Max}$  then return  $\max_{a \in \chi(h)} \text{ExpectiMinimax}(\sigma(h, a), d-1)$

else if  $\rho(h) = \text{Min}$  then return  $\min_{a \in \chi(h)} \text{ExpectiMinimax}(\sigma(h, a), d-1)$

else return  $\sum_{a \in \chi(h)} \text{Pr}[a | h] \text{ExpectiMinimax}(\sigma(h, a), d-1)$

$Z = \{\text{terminal nodes}\}$

$\rho(h)$  = the player to move at  $h$

$\chi(h) = \{\text{available actions at } h\}$

$\sigma(h, a)$  = child of  $h$  produced by  $a$

$\mathbf{u}(h)$  = utility profile for  $h$

$\mathbf{v}[i]$  =  $i$ 'th element of  $\mathbf{v}$


- If  $d \geq \text{height}(h)$ , returns Max's SPE expected payoff

➤ Otherwise returns an approximation

- Can be modified to return an action

- Can also be modified to do  $\alpha$ - $\beta$  pruning

➤ But it's more complicated and less effective than in deterministic games




# In Practice

- Dice rolls increase branching factor
  - 21 possible rolls with 2 dice
  - Given the dice roll,  $\approx 20$  legal moves on average
    - For some dice roles, can be much higher
      - At depth 4, number of nodes is  $= 20 \times (21 \times 20)^3 \approx 1.2 \times 10^9$
  - As depth increases, probability of reaching a given node shrinks
    - $\Rightarrow$  value of lookahead is diminished
- $\alpha$ - $\beta$  pruning is much less effective
- TDGammon (1992) used depth-2 search + very good evaluation function
  - Created its evaluation function automatically using a machine-learning technique called *Temporal Difference* learning
    - hence the *TD* in TDGammon
  - $\approx$  world-champion level

# Generalize

- Finite perfect-information games with
  - Multiple players
  - Chance nodes
  - Nonzero-sum payoffs





# Expectimax

$Z = \{\text{terminal nodes}\}$


$\rho(h) = \text{the player to move at } h$

$\chi(h) = \{\text{available actions at } h\}$

$\sigma(h, a) = \text{child of } h \text{ produced by } a$

$\mathbf{u}(h) = \text{utility profile for } h$

$\mathbf{v}[i] = i\text{'th element of } \mathbf{v}$



function Expectimax( $h, d$ )

if  $h \in Z$  then return  $\mathbf{u}(h)$

else if  $d = 0$  then return  $\mathbf{e}(h)$

if  $h$  is a chance node then return  $\sum_{a \in \chi(h)} \text{Pr}[a | h] \text{ Expectimax}(\sigma(h, a), d-1)$

$V = \{\text{Expectimax}(\sigma(h, a), d-1) \mid a \in \chi(h)\}$

return  $\arg \max_{\mathbf{v} \in V} \mathbf{v}[\rho(h)]$

# Expectimax

function Expectimax( $h, d$ )

if  $h \in Z$  then return  $\mathbf{u}(h)$

else if  $d = 0$  then return  $\mathbf{e}(h)$

if  $h$  is a chance node then return  $\sum_{a \in \chi(h)} \Pr[a \mid h] \text{Expectimax}(\sigma(h, a), d-1)$

$V = \{\text{Expectimax}(\sigma(h, a), d-1) \mid a \in \chi(h)\}$

return  $\arg \max_{\mathbf{v} \in V} \mathbf{v}[\rho(h)]$

function Expectimax-choice( $h, d$ )

if  $h \in Z$  then return *error*

return  $\arg \max_{a \in \chi(h)} (\text{Expectimax}(\sigma(h, a), d-1))[\rho(h)]$

- If  $d \geq \text{height}(h)$  then

- Expectimax returns the SPE payoff profile

- Expectimax-choice returns the SPE action for player  $\rho(h)$

$Z = \{\text{terminal nodes}\}$

$\rho(h) = \text{the player to move at } h$

$\chi(h) = \{\text{available actions at } h\}$

$\sigma(h, a) = \text{child of } h \text{ produced by } a$

$\mathbf{u}(h) = \text{utility profile for } h$

$\mathbf{v}[i] = i\text{'th element of } \mathbf{v}$

# Summary

- Finite two-player zero-sum perfect-information games
  - $\text{maxmin} = \text{minmax} = \text{Nash}$
  - only need to look at pure strategies
  - game-tree search
    - Minimax, LD-minimax, Alpha-Beta
    - limited search depth, static evaluation function
  - Monte Carlo roll-outs, UCT
- Multiplayer games
  - $\text{Max}^n$  algorithm (depth-limited backward induction)
  - Paranoid algorithm (approximate maxmin)
- Games with chance nodes
  - Expectiminimax (2-player zero-sum)
  - Expectimax ( $n$ -player nonzero-sum)